



UNSW
SYDNEY

Relational Algebra

COMP9311 25T3; Week 2.2

By Xubo Wang, UNSW

Motivation

- We've seen what a relational model is.
- We needed a *formal* language to specify data (tuples) from the relational model.
- **Relational Algebra** (E.F. Codd (1970))

Why Relational Algebra?

- It provides a formal foundation for relational model operations
- It is used as a basis for implementing and optimizing queries in the query processing and optimization
- Some of its concepts are incorporated into SQL

Relational Algebra

Relational Algebra is a procedural data manipulation language (**DML**).

It specifies operations on relations to define new relations:

Unary Relational Operations: Select, Project

Operations from Set Theory: Union, Intersection, Difference,
Cartesian Product

Binary Relational Operations: Join, Divide.

1 SELECT

The SELECT operation/predicate is used to select a subset of the tuples of a relation R, satisfying some conditions.

Notation: $\sigma_{<selection\ condition>} (R)$

Intuition: Filters out all tuples that do not satisfy select condition



Selection Condition

The condition is defined by a ***selection clause***:

- <attribute> *operator* <constant>
- <attribute> *operator* <attribute>

Where *operator* is one of =, <, ≤, >, ≥ or ≠ ...

Example:

- age ≤ 24
- commission ≥ 24 000

Selection Condition

Selection clauses can also be

- <expression> operator <expression>

With this, we can use **Boolean connectives** as operators

- C1 AND C2
- C1 OR C2
- NOT C

Terms equivalently expressed by \wedge (and), \vee (or), \neg (not)

Q: Select the enrolment records for the students whose supervisor is Person 1

ENROLMENT:

Enrolment#	Supervisee	Supervisor	Department	Degree
1	1	2	Psychology	Ph.D.
2	3	1	Comp.Sci.	Ph.D.
3	4	1	Comp.Sci.	M.Sc.
4	5	1	Comp.Sci.	M.Sc.

$$\sigma_{(Supervisor=1)}(ENROLMENT)$$

The output relation is

Enrolment#	Supervisee	Supervisor	Department	Degree
2	3	1	Comp.Sci	Ph.D.
3	4	1	Comp.Sci	M.Sc.
4	5	1	Comp.Sci	M.Sc.

Q: Select the enrolment records for Person 1's non-Ph.D. students

ENROLMENT:

Enrolment#	Supervisee	Supervisor	Department	Degree
1	1	2	Psychology	Ph.D.
2	3	1	Comp.Sci.	Ph.D.
3	4	1	Comp.Sci.	M.Sc.
4	5	1	Comp.Sci.	M.Sc.

$$\left. \begin{array}{l} \sigma_{(Supervisor=1 \text{ AND } Degree \neq "Ph.D.")}(ENROLMENT) \\ \sigma_{(Supervisor=1 \text{ AND } NOT Degree = "Ph.D.")}(ENROLMENT) \end{array} \right\} \text{Same}$$

The output relation is

Enrolment#	Supervisee	Supervisor	Department	Degree
3	4	1	Comp.Sci	M.Sc.
4	5	1	Comp.Sci	M.Sc.

Properties of Selection

Properties:

- Consecutive selects ***can be combined***:

$$\sigma_{<cond_1>}(\sigma_{<cond_2>}(R)) = \sigma_{<cond_1> \text{ AND } <cond_2>}(R)$$

- Selection is a ***commutative*** operation:

$$\sigma_{<cond_1>}(\sigma_{<cond_2>}(R)) = \sigma_{<cond_2>}(\sigma_{<cond_1>}(R))$$

2 PROJECT

The PROJECT operation is used to project a subset of the attributes (columns) of a relation, denoted by:

General form: $\pi_{<attribute\ list>}(R)$

Result:

- schema: attribute list (A_1, \dots, A_k)
- instance: the set of all subtuples $t[A_1, \dots, A_k]$ where $t \in R$

Q: Find departments and degree requirements for the courses that students enroll.

ENROLMENT:

Enrolment#	Supervisee	Supervisor	Department	Degree
1	1	2	Psychology	Ph.D.
2	3	1	Comp.Sci.	Ph.D.
3	4	1	Comp.Sci.	M.Sc.
4	5	1	Comp.Sci.	M.Sc.

$$\pi_{\{department, degree\}}(ENROLMENT)$$

The output relation is

Department	Degree
Psychology	Ph.D.
Comp.Sci	Ph.D.
Comp.Sci	M.Sc.

Duplicates of PROJECT

ENROLMENT:

Enrolment#	Supervisee	Supervisor	Department	Degree
1	1	2	Psychology	Ph.D.
2	3	1	Comp.Sci.	Ph.D.
3	4	1	Comp.Sci.	M.Sc.
4	5	1	Comp.Sci.	M.Sc.

Question: What if we do PROJECTION on only department?

Department
Psychology
Comp.Sci.
Comp.Sci.
Comp.Sci.

or

Department
Psychology
Comp.Sci.

?

Duplicates of PROJECT

ENROLMENT:

Enrolment#	Supervisee	Supervisor	Department	Degree
1	1	2	Psychology	Ph.D.
2	3	1	Comp.Sci.	Ph.D.
3	4	1	Comp.Sci.	M.Sc.
4	5	1	Comp.Sci.	M.Sc.

Question: What if we do PROJECTION on only department?

Answer: Keep only one ‘Comp.Sci.’.

Department
Psychology
Comp.Sci.

Relational Algebra is based on sets, so no duplicates are allowed.

- The PROJECT operation *removes any duplicate tuples*, so the result of the PROJECT operation is a set of distinct tuples, and this is known as **duplicate elimination**.

Properties of PROJECT

Consider $\pi_{\langle \text{list1} \rangle}(\pi_{\langle \text{list2} \rangle}(R))$

If <list2> contains all the attributes in <list1>:

Then $\pi_{\langle \text{list1} \rangle}(\pi_{\langle \text{list2} \rangle}(R)) = \pi_{\langle \text{list1} \rangle}(R)$

Else the operation is *not well defined.*

Project Predicate

Question: is projection commutative with selection?

i.e., $\pi_X(\sigma_B(R)) = \sigma_B(\pi_X(R))$?

Consider the following:

$$\pi_{\{degree\}}(\sigma_{(Department='Psychology')}(ENROLMENT))$$

Degree
Ph.D.

$$\sigma_{(Department='Psychology')}(\pi_{\{degree\}}(ENROLMENT))$$

Error as SELECT cannot find Department

Answer: The attribute used in SELECT must be a subset of the attribute list in PROJECT

Project Predicate

Question: is projection commutative with selection?

i.e., $\pi_X(\sigma_B(R)) = \sigma_B(\pi_X(R))$?

Consider the following:

$$\pi_{\{Department, Degree\}}(\sigma_{(Department='Psychology')}(ENROLMENT))$$
$$\sigma_{(Department='Psychology')}(\pi_{\{Department, Degree\}}(ENROLMENT))$$

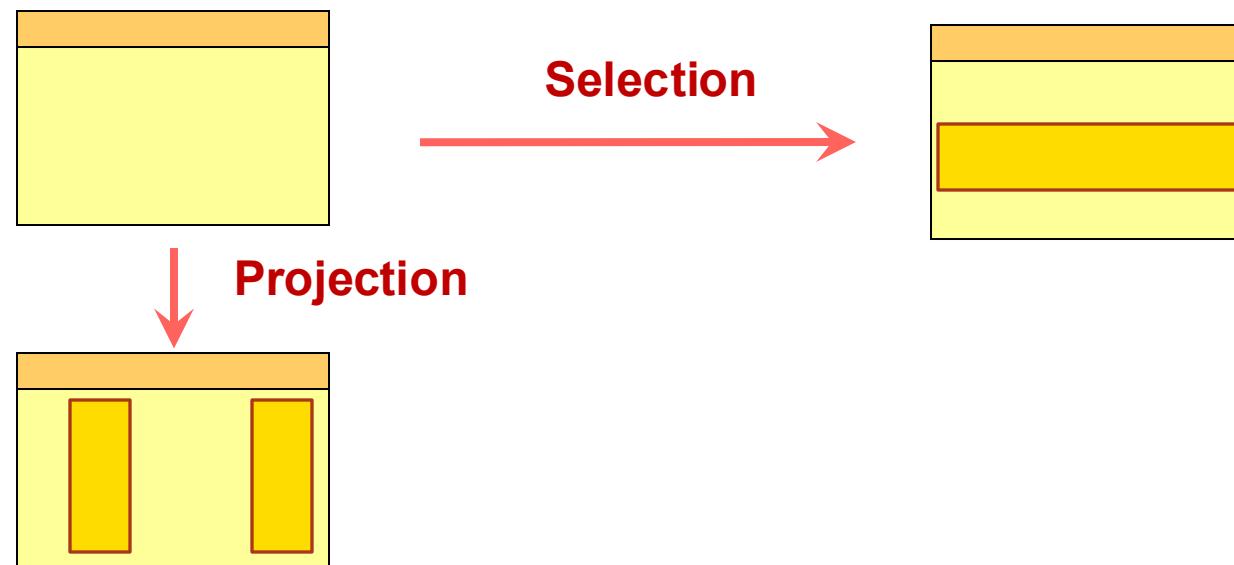
Enrolment#	Supervisee	Supervisor	Department	Degree
1	1	2	Psychology	Ph.D.
2	3	1	Comp.Sci.	Ph.D.
3	4	1	Comp.Sci.	M.Sc.
4	5	1	Comp.Sci.	M.Sc.

Department	Degree
Psychology	Ph.D.

Answer: The attribute used in SELECT must be a subset of the attribute list in PROJECT

Intuition: Projection and Selection

1. Selection performs a horizontal decomposition, and
2. Projection performs a vertical decomposition



3 SET UNION

UNION is the set-theoretic union of the tuples of two relations.

$$R \cup S = \{t: t \in R \text{ or } t \in S\}$$

Condition: R and S must be **union compatible!**

Union compatibility: there is a 1-1 correspondence between their attributes:
the same **name** and same **domain**.

Example: `Section(course_id, semester, year, instructor)`

To find all courses taught in the Fall 2009 semester, **or** in the Spring 2010 semester, **or** in both:

$$\begin{aligned} & \pi_{\{course_id\}}(\sigma_{(semester="Fall" \wedge year=2009)}(section)) \cup \\ & \pi_{\{course_id\}}(\sigma_{(semester="Spring" \wedge year=2010)}(section)) \end{aligned}$$

Example

STUDENT:

Person#	Name
1	Dr C.C.Chen
3	Ms K.Juliff
4	Ms J.Gledill
5	Ms B.K.Lee

RESEARCHER:

Person#	Name
1	Dr C.C.Chen
2	Dr R.G.Wilkinson

Example: $STUDENT \cup RESEARCHER =$

Person#	Name
1	Dr C.C.Chen
3	Ms K.Juliff
4	Ms J.Gledhill
5	Ms B.K.Lee
2	Dr R.G.Wilkinson

4 SET INTERSECTION

Enrolment#	Supervisee	Supervisor	Department	Degree
1	1	2	Psychology	Ph.D.
2	3	1	Comp.Sci.	Ph.D.
3	4	1	Comp.Sci.	M.Sc.
4	5	1	Comp.Sci.	M.Sc.

- *INTERSECTION* is an operation that includes all tuples that are present in **both** relations, denoted by

$$R \cap S = \{t: t \in R \text{ and } t \in S\}$$

- Condition: R and S must also be **union compatible!**
- Example: $R_1 \leftarrow \sigma_{(supervisor=1)}(ENROLMENT)$
 $R_2 \leftarrow \sigma_{(degree='Ph.D.')}(ENROLMENT)$

$$R_1 \cap R_2 =$$

Enrolment#	Supervisee	Supervisor	Department	Degree
2	3	1	Comp.Sci.	Ph.D.

Example of Intersection

STUDENT:

Person#	Name
1	Dr C.C.Chen
3	Ms K.Juliff
4	Ms J.Gledill
5	Ms B.K.Lee

RESEARCHER:

Person#	Name
1	Dr C.C.Chen
2	Dr R.G.Wilkinson

Example: STUDENT \cap RESEARCHER =

Person#	Name
1	Dr C.C. Chen

5 SET DIFFERENCE

DIFFERENCE is an operation that includes all tuples that are in the left relation but not in the right relation, denoted by

$$R - S = \{t: t \in R \text{ and } t \notin S\}$$

Condition: R and S must also be **union compatible!**

STUDENT:

Person#	Name
1	Dr C.C.Chen
3	Ms K.Juliff
4	Ms J.Gledill
5	Ms B.K.Lee

RESEARCHER:

Person#	Name
1	Dr C.C.Chen
2	Dr R.G.Wilkinson

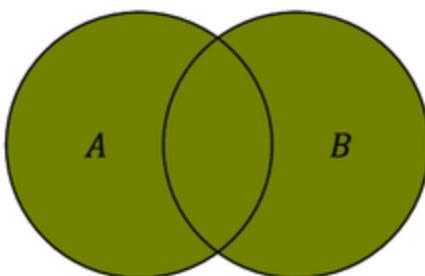
Example: STUDENT – RESEARCHRER =

Person#	Name
3	Ms K. Juliff
4	Ms J. Gledhill
5	Ms B.K. Lee

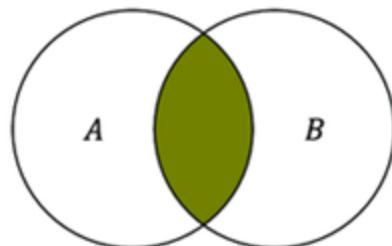
Summary

Operations on Relations

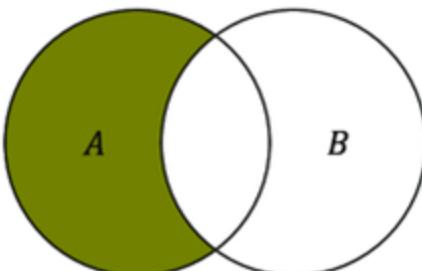
- *UNION:* $A \cup B$



- *INTERSECTION:* $A \cap B$



- *DIFFERENCE:* $A - B$



Express: The names of persons who are either a student or a researcher

STUDENT:

Person#	Name
1	Dr C.C.Chen
3	Ms K.Juliff
4	Ms J.Gledill
5	Ms B.K.Lee

RESEARCHER:

Person#	Name
1	Dr C.C.Chen
2	Dr R.G.Wilkinson

Express: The names of persons who are either a student or a researcher

STUDENT:

Person#	Name
1	Dr C.C.Chen
3	Ms K.Juliff
4	Ms J.Gledill
5	Ms B.K.Lee

RESEARCHER:

Person#	Name
1	Dr C.C.Chen
2	Dr R.G.Wilkinson

$$\pi_{\{name\}}(STUDENT \cup RESEARCHER)$$

Name
Dr C.C.Chen
Dr
R.G.Wilkinson
Ms K.Juliff
Ms J.Gledill
Ms B.K.Lee

Express: The names of persons who are a student **and** a researcher

STUDENT:

Person#	Name
1	Dr C.C.Chen
3	Ms K.Juliff
4	Ms J.Gledill
5	Ms B.K.Lee

RESEARCHER:

Person#	Name
1	Dr C.C.Chen
2	Dr R.G.Wilkinson

Express: The names of persons who are a student **and** a researcher

STUDENT:

Person#	Name
1	Dr C.C.Chen
3	Ms K.Juliff
4	Ms J.Gledill
5	Ms B.K.Lee

RESEARCHER:

Person#	Name
1	Dr C.C.Chen
2	Dr R.G.Wilkinson

$$\pi_{\{name\}}(STUDENT \cap RESEARCHER)$$

Name
Dr C.C.Chen

Express: The names of persons who are a student but not a researcher

STUDENT:

<u>Person#</u>	<u>Name</u>
1	Dr C.C.Chen
3	Ms K.Juliff
4	Ms J.Gledill
5	Ms B.K.Lee

RESEARCHER:

<u>Person#</u>	<u>Name</u>
1	Dr C.C.Chen
2	Dr R.G.Wilkinson

Express: The names of persons who are a student but not a researcher

STUDENT:

Person#	Name
1	Dr C.C.Chen
3	Ms K.Juliff
4	Ms J.Gledill
5	Ms B.K.Lee

RESEARCHER:

Person#	Name
1	Dr C.C.Chen
2	Dr R.G.Wilkinson

$$\pi_{\{name\}}(STUDENT - RESEARCHER)$$

Name
Ms K.Juliff
Ms J.Gledill
Ms B.K.Lee

Express: The departments and degrees of Courses which are not enrolled by any student

ENROLMENT:

<u>Enrolment#</u>	<u>Supervisee</u>	<u>Supervisor</u>	<u>Department</u>	<u>Degree</u>
1	1	2	Psychology	Ph.D.
2	3	1	Comp.Sci.	Ph.D.
3	4	1	Comp.Sci.	M.Sc.
4	5	1	Comp.Sci.	M.Sc.

COURSE:

<u>Department</u>	<u>Degree</u>
Psychology	Ph.D.
Comp.Sci.	Ph.D.
Comp.Sci.	M.Sc.
Psychology	M.Sc.

Express: The departments and degrees of Courses which are not enrolled by any student

ENROLMENT:

<u>Enrolment#</u>	<u>Supervisee</u>	<u>Supervisor</u>	<u>Department</u>	<u>Degree</u>
1	1	2	Psychology	Ph.D.
2	3	1	Comp.Sci.	Ph.D.
3	4	1	Comp.Sci.	M.Sc.
4	5	1	Comp.Sci.	M.Sc.

COURSE:

<u>Department</u>	<u>Degree</u>
Psychology	Ph.D.
Comp.Sci.	Ph.D.
Comp.Sci.	M.Sc.
Psychology	M.Sc.

COURSE – ($\pi_{(Department, Degree)}(ENROLMENT)$)

CARTESIAN PRODUCT

$$R \times S = \{t_1 \parallel t_2 : t_1 \in R \text{ and } t_2 \in S\}$$

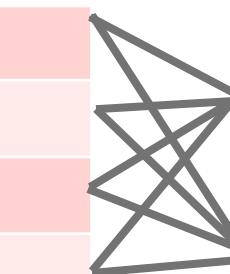
- Intuition: **every combination of tuples in R with tuples in S.**
- $t_1 \parallel t_2$ indicates the concatenation of tuples.
- R and S not required to be union compatible, but
- The number of tuples in the output relations is always $|R| * |S|$

Usually assumes that attributes of $r(R)$ and $s(S)$ are disjoint. (That is, $R \cap S = \emptyset$). If not, you must devise a naming schema to distinguish between the attribute names if they are the same for example attribute A in $r(A, B)$ and $s(A, C)$, by attaching the relation's name, $r.A$ and $s.A$ (known as **dot-notation**)

Example of cartesian product

ENROLMENT:

<u>Enrolment#</u>	Supervisee	Supervisor	Department	Degree
1	1	2	Psychology	Ph.D.
2	3	1	Comp.Sci.	Ph.D.
3	4	1	Comp.Sci.	M.Sc.
4	5	1	Comp.Sci.	M.Sc.



RESEARCHER:

<u>Person#</u>	Name
1	Dr C.C.Chen
2	Dr R.G.Wilkinson

Person#	Name
1	Dr C.C.Chen
2	Dr R.G.Wilkinson

Example of cartesian product

ENROLMENT X RESEARCHER =

Enrolment#	Supervisee	Supervisor	Department	Degree
1	1	2	Psychology	Ph.D.
2	3	1	Comp.Sci.	Ph.D.
3	4	1	Comp.Sci.	M.Sc.
4	5	1	Comp.Sci.	M.Sc.

E'ment#	S'ee	S'or	D'ment	Degree	Person#	Name
1	1	2	Psych.	Ph.D.	1	Dr C.C. Chen
1	1	2	Psych.	Ph.D.	2	Dr R.G. Wilkinson
2	3	1	Cmp.Sci	Ph.D.	1	Dr C.C. Chen
2	3	1	Cmp.Sci	Ph.D.	2	Dr R.G. Wilkinson
3	4	1	Cmp.Sci	M.Sc.	1	Dr C.C. Chen
3	4	1	Cmp.Sci	M.Sc.	2	Dr R.G. Wilkinson
4	5	1	Cmp.Sci	M.Sc.	1	Dr C.C. Chen
4	5	1	Cmp.Sci	M.Sc.	2	Dr R.G. Wilkinson

There were 4 tuples in ENROLMENT and 2 tuples in RESEARCHER. In the result, there are 8 tuples.

Useful if we add a condition

$$R_1 \leftarrow ENROLMENT \times RESEARCHER$$

E'ment#	S'ee	S'or	D'ment	Degree	Person#	Name
1	1	2	Psych.	Ph.D.	1	Dr C.C. Chen
1	1	2	Psych.	Ph.D.	2	Dr R.G. Wilkinson
2	3	1	Cmp.Sci	Ph.D.	1	Dr C.C. Chen
2	3	1	Cmp.Sci	Ph.D.	2	Dr R.G. Wilkinson
3	4	1	Cmp.Sci	M.Sc.	1	Dr C.C. Chen
3	4	1	Cmp.Sci	M.Sc.	2	Dr R.G. Wilkinson
4	5	1	Cmp.Sci	M.Sc.	1	Dr C.C. Chen
4	5	1	Cmp.Sci	M.Sc.	2	Dr R.G. Wilkinson

In practice it's useful if we give a cartesian product specified condition

$$\sigma_{(Supervisor=Person\#)}(R_1) =$$

E'ment#	S'ee	S'or	D'ment	Degree	Person#	R'cher. Name
1	1	2	Psych.	Ph.D.	2	Dr R.G. Wilkinson
2	3	1	Cmp.Sci.	Ph.D.	1	Dr C.C. Chen
3	4	1	Cmp.Sci.	M.Sc.	1	Dr C.C. Chen
4	5	1	Cmp.Sci.	M.Sc.	1	Dr C.C. Chen

More useful if we add a projection

$R_1 \leftarrow ENROLMENT \times RESEARCHER$

$R_2 \leftarrow \sigma_{(Supervisor=Person\#)}(R_1)$

E'ment#	S'ee	S'or	D'ment	Degree	Person#	R'cher. Name
1	1	2	Psych.	Ph.D.	2	Dr R.G.Wilkinson
2	3	1	Cmp.Sci.	Ph.D.	1	Dr C.C. Chen
3	4	1	Cmp.Sci.	M.Sc.	1	Dr C.C. Chen
4	5	1	Cmp.Sci.	M.Sc.	1	Dr C.C. Chen

$\pi_{\{E'ment\#, S'ee, S'or, Name, D'ment, Degree\}}(R_2)$

E'ment#	S'ee	S'or	Name	D'ment	Degree
1	1	2	Dr R.G.Wilkinson	Psych.	Ph.D.
2	3	1	Dr C.C. Chen	Comp.Sci.	Ph.D.
3	4	1	Dr C.C. Chen	Comp.Sci.	M.Sc.
4	5	1	Dr C.C. Chen	Comp.Sci.	M.Sc.

The two equal attributes occur only once

The last of these is also known as *natural join*, the next to last is *equi-join*.

6 JOIN

- JOIN is used to combine related tuples from two relations into single "longer" tuples.
- **Theta-join**

$$R \bowtie_{\text{join condition}} S = \{t_1 \parallel t_2 : t_1 \in R \text{ and } t_2 \in S \text{ and } \text{join condition}\}$$

- A general join condition is of the form:
<condition> AND <condition> AND ... AND <condition>

<condition> is of the form $A_i \theta B_j$, A_i is an attribute of R , B_j is an attribute of S , A_i and B_j have the same domain, and θ (theta) is one of the comparison operators $\{=, <, \leq, >, \geq, \neq\}$.

6.1 Equi-join

A type of theta-join where the only comparison operator used is “=” is called an Equi-join

Example:

$$ENROLMENT \bowtie_{(Supervisor=Person\#)} RESEARCHER$$

$$R_2 \leftarrow \sigma_{(Supervisor=Person\#)}(R_1)$$

E'ment#	S'ee	S'or	D'ment	Degree	Person#	R'cher. Name
1	1	2	Psych.	Ph.D.	2	Dr R.G. Wilkinson
2	3	1	Cmp.Sci.	Ph.D.	1	Dr C.C. Chen
3	4	1	Cmp.Sci.	M.Sc.	1	Dr C.C. Chen
4	5	1	Cmp.Sci.	M.Sc.	1	Dr C.C. Chen

6.2 Natural Join

ENROLMENT:					COURSE:	
Enrolment#	Supervisee	Supervisor	Department	Degree	Department	Degree
1	1	2	Psychology	Ph.D.	Psychology	Ph.D.
2	3	1	Comp.Sci.	Ph.D.	Comp.Sci.	Ph.D.
3	4	1	Comp.Sci.	M.Sc.	Comp.Sci.	M.Sc.
4	5	1	Comp.Sci.	M.Sc.	Psychology	M.Sc.

A type of equi-join that requires each pair of join attributes to have the same name and domain in both relations.

Notes: In a natural join, there may be several valid pairs of join attributes.

$ENROLMENT \bowtie_{(Department, Degree), (Deparment, Degree)} COURSE$

If there are pairs of joining attributes identically named, we can write

$ENROLMENT \bowtie COURSE$

Note: this notion also acceptable if there's one join attribute

6.2 Natural Join

Intuitions:

- Enforce equality on all attributes with same name
- Eliminate one copy of duplicated attributes

E'ment#	S'ee	S'or	Name	D'ment	Degree
1	1	2	Dr R.G.Wilkinson	Psych.	Ph.D.
2	3	1	Dr C.C. Chen	Comp.Sci.	Ph.D.
3	4	1	Dr C.C. Chen	Comp.Sci.	M.Sc.
4	5	1	Dr C.C. Chen	Comp.Sci.	M.Sc.

JOINS

Remember the differences between the types of joins:

1. Theta JOIN

$$R \bowtie_{\text{join condition}} S$$

2. Equi JOIN

3. Natural JOIN

Note: all denoted with \bowtie

Pracs

STUDENT:

<u>Person#</u>	Name
1	Mr J.He
3	Ms K.Juliff
4	Ms J.Gledill
5	Ms B.K.Lee

RESEARCHER:

<u>Person#</u>	Name
1	Dr C.C.Chen
2	Dr R.G.Wilkinson

COURSE

<u>Depart</u>	<u>Degree</u>
EE	PhD
CS	PhD
EE	MSc
CS	MSc

ENROLMENT:

<u>Enrol#</u>	Supervisee	Supervisor	Depart	Degree
1	1	2	EE	PhD
2	3	1	CS	PhD
3	4	1	CS	MSc
4	5	1	CS	MSc

What are the names of students who are studying for an MSc in CS?

Pracs

STUDENT:

<u>Person#</u>	Name
1	Mr J.He
3	Ms K.Juliff
4	Ms J.Gledill
5	Ms B.K.Lee

RESEARCHER:

<u>Person#</u>	Name
1	Dr C.C.Chen
2	Dr R.G.Wilkinson

COURSE

<u>Depart</u>	<u>Degree</u>
EE	PhD
CS	PhD
EE	MSc
CS	MSc

ENROLMENT:

<u>Enrol#</u>	Supervisee	Supervisor	Depart	Degree
1	1	2	EE	PhD
2	3	1	CS	PhD
3	4	1	CS	MSc
4	5	1	CS	MSc

What are the names of students who are studying for an MSc in CS?

$\pi_{\{name\}}(\sigma_{(degree=MSc \text{ and } Depart=CS)} ENROLMENT \bowtie_{supervisee=person\#} Student)$

Pracs

STUDENT:

<u>Person#</u>	Name
1	Mr J.He
3	Ms K.Juliff
4	Ms J.Gledill
5	Ms B.K.Lee

RESEARCHER:

<u>Person#</u>	Name
1	Dr C.C.Chen
2	Dr R.G.Wilkinson

COURSE

<u>Depart</u>	<u>Degree</u>
EE	PhD
CS	PhD
EE	MSc
CS	MSc

ENROLMENT:

<u>Enrol#</u>	Supervisee	Supervisor	Depart	Degree
1	1	2	EE	PhD
2	3	1	CS	PhD
3	4	1	CS	MSc
4	5	1	CS	MSc

The IDs of students who are supervised by Dr C.C.Chen

Pracs

STUDENT:

<u>Person#</u>	Name
1	Mr J.He
3	Ms K.Juliff
4	Ms J.Gledill
5	Ms B.K.Lee

RESEARCHER:

<u>Person#</u>	Name
1	Dr C.C.Chen
2	Dr R.G.Wilkinson

COURSE

<u>Depart</u>	<u>Degree</u>
EE	PhD
CS	PhD
EE	MSc
CS	MSc

ENROLMENT:

<u>Enrol#</u>	Supervisee	Supervisor	Depart	Degree
1	1	2	EE	PhD
2	3	1	CS	PhD
3	4	1	CS	MSc
4	5	1	CS	MSc

The IDs of students who are supervised by Dr C.C.Chen

R1 = ENROLMENT \bowtie (supervisor=person#) RESEARCHER

R2 = $\sigma_{(name=Dr\ C.C.Chen)}$ R1

R3 = $\pi_{\{supervisee\}}$ R2

Divide

- The DIVISION operation is applied to two Relations R and S, where the attributes of S are a subset of the attributes of R.
- The relation returned by the division operator will have attributes = (All attributes of R – All Attributes of S)
- Return all tuples from relation R which are associated to every S's tuple.

R $R \div S =$

A	B
a ₁	b ₁
a ₁	b ₂
a ₂	b ₁
a ₃	b ₂
a ₄	b ₁
a ₅	b ₁
a ₅	b ₂

S

B
b ₁
b ₂

A

A
a ₁
a ₅

Divide

Typical use: which departments offer all degrees?

$$Course \div (\pi_{Degree} Course)$$

COURSE

<u>Depart</u>	<u>Degree</u>
EE	PhD
CS	PhD
EE	MSc
CS	MSc

Pracs

STUDENT:

<u>Person#</u>	Name
1	Mr J.He
3	Ms K.Juliff
4	Ms J.Gledill
5	Ms B.K.Lee

RESEARCHER:

<u>Person#</u>	Name
1	Dr C.C.Chen
2	Dr R.G.Wilkinson

COURSE

<u>Depart</u>	<u>Degree</u>
EE	PhD
CS	PhD
EE	MSc
CS	MSc

ENROLMENT:

<u>Enrol#</u>	Supervisee	Supervisor	Depart	Degree
1	1	2	EE	PhD
2	3	1	CS	PhD
3	4	1	CS	MSc
4	5	1	CS	MSc

The names of supervisor who supervises both MSc and PhD students

Pracs

STUDENT:

<u>Person#</u>	Name
1	Mr J.He
3	Ms K.Juliff
4	Ms J.Gledill
5	Ms B.K.Lee

RESEARCHER:

<u>Person#</u>	Name
1	Dr C.C.Chen
2	Dr R.G.Wilkinson

COURSE

<u>Depart</u>	<u>Degree</u>
EE	PhD
CS	PhD
EE	MSc
CS	MSc

ENROLMENT:

<u>Enrol#</u>	Supervisee	Supervisor	Depart	Degree
1	1	2	EE	PhD
2	3	1	CS	PhD
3	4	1	CS	MSc
4	5	1	CS	MSc

The names of supervisor who supervises both MSc and PhD students

R1 = $\pi_{\{\text{SUPERVISOR}, \text{DEGREE}\}} \text{ ENROLMET} \div \pi_{\{\text{DEGREE}\}} \text{ COURSE}$

R2 = $\pi_{\{\text{Name}\}} (\text{R1} \bowtie_{(\text{supervisor}=\text{person}\#)} \text{ RESEARCH})$

Exercise

R:

A	B	C
a ₁	b ₁	c ₁
a ₁	b ₁	c ₂
a ₁	b ₁	c ₃
a ₁	b ₂	c ₂
a ₂	b ₁	c ₁
a ₂	b ₂	c ₂
a ₃	b ₁	c ₁
a ₃	b ₂	c ₁
a ₃	b ₂	c ₂

S:

B	C
b ₁	c ₁
b ₁	c ₂

Write relational algebra that retrieves:

1. Find A of R that contains all S.

Exercise

R:

A	B	C
a ₁	b ₁	c ₁
a ₁	b ₁	c ₂
a ₁	b ₁	c ₃
a ₁	b ₂	c ₂
a ₂	b ₁	c ₁
a ₂	b ₂	c ₂
a ₃	b ₁	c ₁
a ₃	b ₂	c ₁
a ₃	b ₂	c ₂

S:

B	C
b ₁	c ₁
b ₁	c ₂

Write relational algebra that retrieves:

1. Find A of R that contains all S.

$R \div S$

A

a₁

Exercise

R:

A	B	C
a ₁	b ₁	c ₁
a ₁	b ₁	c ₂
a ₁	b ₁	c ₃
a ₁	b ₂	c ₂
a ₂	b ₁	c ₁
a ₂	b ₂	c ₂
a ₃	b ₁	c ₁
a ₃	b ₂	c ₁
a ₃	b ₂	c ₂

S:

B	C
b ₁	c ₁
b ₁	c ₂

Write relational algebra that retrieves:

2. Find (A, B) of R that contains all C of S.

Exercise

R:

A	B	C
a ₁	b ₁	c ₁
a ₁	b ₁	c ₂
a ₁	b ₁	c ₃
a ₁	b ₂	c ₂
a ₂	b ₁	c ₁
a ₂	b ₂	c ₂
a ₃	b ₁	c ₁
a ₃	b ₂	c ₁
a ₃	b ₂	c ₂

S:

B	C
b ₁	c ₁
b ₁	c ₂

Write relational algebra that retrieves:

2. Find (A, B) of R that contains all C of S.

$$R \div \pi_{\{C\}}(S)$$

A	B
a ₁	b ₁
a ₃	b ₂

Rename Operator

- The **rename** operator ρ changes the name of one or more attributes
- Change the names in a schema
- Does not affect **instance** of the target relation

Family

Father	Child
Adam	Abel
Adam	Cain
Abraham	Isaac

$\rho_{(\text{Parent}, \text{Child})}(\text{Family})$

Parent	Child
Adam	Abel
Adam	Cain
Abraham	Isaac

$$\rho_{S(A_1, A_2, \dots, A_n)}(R) \quad \text{or} \quad \rho_S(R) \quad \text{or} \quad \rho_{(A_1, A_2, \dots, A_n)}(R)$$

- Why might this be useful? To be included in relational algebra?

Why RENAME Operator?

- To unify schemas for set operators
 - **Union compatibility:** there is a 1-1 correspondence between their attributes: the same **name** and same **domain**.

$$R(zid, name), S(studentid, name)$$
$$R \cup \rho_{(zid, name)} S$$

- For disambiguation in “self-join”

$$STUDENT \bowtie_{STUDENT.major = STUDENT.major} STUDENT$$
$$\rho_{S1}(STUDENT) \bowtie_{S1.major = S2.major} \rho_{S2}(STUDENT)$$

Basic vs Extended Operators

Note: $\{\sigma, \pi, U, -, \times\}$ (and *rename*) are sufficient to define all these operations: this is a relationally complete set of operators. These are the **basic operators** of the Relational Algebra.

What about *JOIN*, *INTERSECTION* and *DIVIDE*?

They are **extended operators** because they can be derived from the basic operators.

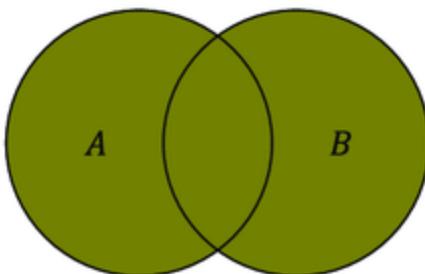
$$R \bowtie_{< \text{condition} >} S = \sigma_{< \text{condition} >} (R \times S)$$

$$A \cap B = A - (A - B)$$

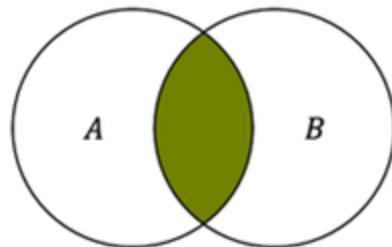
Summary

Operations on Relations

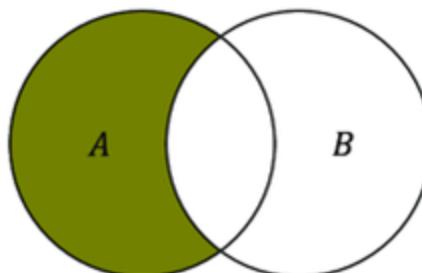
- *UNION:* $A \cup B$



- *INTERSECTION:* $A \cap B$



- *DIFFERENCE:* $A - B$



$$A \cap B = A - (A - B)$$

Basic vs Extended Operators

Note: $\{\sigma, \pi, U, -, \times\}$ (and *rename*) are sufficient to define all these operations: this is a relationally complete set of operators. These are the **basic operators** of the Relational Algebra.

What about *JOIN*, *INTERSECTION* and *DIVIDE*?

They are **extended operators** because they can be derived from the basic operators.

E.g., We can write $R \div S$ as

$$TEMP1 \leftarrow \pi_{R-S}(R)$$

$$TEMP2 \leftarrow \pi_{R-S}((TEMP1 \times S) - R)$$

$$RESULT = TEMP1 - TEMP2$$

- $R-S$: attributes in R not in S
- The result to the right of \leftarrow is assigned to the relation variable on the left of \leftarrow .
- May use variable in subsequent expressions.

Aggregate Operators

What if we want a relation with information about “sum of salaries” of employees, or the “average age” of students?

We need more expressive power. We can use **aggregation functions** to summarize information from multiple tuples into **aggregate values**.

We can use an **aggregation operator** γ and a function such as *SUM*, *AVG*, *MIN*, *MAX*, or *COUNT*.

In general, duplicates are not considered in the aggregation.

If $R =$

A	B
1	2
3	4
3	5
1	1

, then $\gamma_{\text{SUM}(A)}(R) =$

SUM(A)
8

not considered in

and $\gamma_{\text{AVG}(B)}(R) =$

AVG(B)
3

Aggregate Operators

We can also retrieve aggregate values for groups, like the “sum of employee salaries” *per department* or the “average student age” *per faculty*.

We give γ additional arguments to specify that the aggregation behavior should be based on *groups* (defined by a set of attributes).

If $R =$

a	b
1	2
3	4
3	5
1	3

, then $\gamma_{a, \text{SUM}(b)}(R) =$

a	SUM(b)
1	5
3	9

Formal Definition

A **basic relational algebra expression** is one of the following:

- A relation in the database
- (could also be a) constant relation
 - (fixed set of tuples, e.g., $\{(1,2), (1,3), (2,3)\}$)

A **general relational algebra expression** is constructed out of smaller subexpressions.

Let E_1 and E_2 be relational algebra expressions; the following are all relational-algebra expressions:

- $E_1 \cup E_2$
- $E_1 - E_2$
- $E_1 \times E_2$
- $\sigma_P(E_1)$ where P is predicate on attributes in E_1
- $\pi_S(E_1)$ where S is a set of attributes in E_1
- $\rho_X(E_1)$ where X is the new name for the result of E_1

OPERATION	PURPOSE	NOTATION
SELECT	Selects all tuples that satisfy the selection condition from a relation R	$\sigma_{\langle \text{selection condition} \rangle}(R)$
PROJECT	Produces a new relation with only some of the attributes of R and removes duplicate tuples.	$\pi_{\langle \text{attribute list} \rangle}(R)$
THETA-JOIN	Produces all combinations of tuples from R and S that satisfy the join condition.	$R \bowtie_{\langle \text{join condition} \rangle} S$
EQUI-JOIN	Produces all the combinations of tuples from R and S that satisfy a join condition with only equality comparisons.	$R \bowtie_{\langle \text{join condition} \rangle} S$
NATURAL-JOIN	Same as EQUIJOIN except that the join attributes of S are not included in the resulting relation; if the join attributes have the same names, they do not have to be specified at all.	$R \bowtie_{\langle \text{join condition} \rangle} S$
UNION	Produces a relation that includes all the tuples in R or S or both R and S; R and S must be union compatible.	$R \cup S$
INTERSECTION	Produces a relation that includes all the tuples in both R and S; R and S must be union compatible.	$R \cap S$
DIFFERENCE	Produces a relation that includes all the tuples in R that are not in S; R and S must be union compatible.	$R - S$
CARTESIAN PRODUCT	Produces a relation that has the attributes of R and S and includes as tuples all possible combinations of tuples from R and S.	$R \times S$
DIVISION	Produces a relation T(X) that includes all tuples $t[X]$ in $R(Z)$ that appear in R in combination with every tuple from $S(Y)$, where $Z = X \cup Y$.	$R(Z) \div S(Y)$