

DTS207TC Database Development and Design

Lab 9 Tutorial: Data Storage

Instructor: Xiaowu Sun
School of AI and Advanced Computing (AIAC), XJTLU

November 2025

Contents

1 Lab 9-Part 1: Fixed-Length Record	2
Theoretical Background	2
Example Schema	2
Record Size and Page Capacity	2
Code Implementation	2
Exercise 1.1 — Implement <code>insert()</code> (Dense Prefix)	3
Exercise 1.2 — Implement <code>delete_swap()</code>	4
Exercise 1.3 — Implement <code>delete_compact()</code>	4
Exercise 1.4 — Run and Observe (Practical Demo)	5
2 Lab 9-Part 2: Variable-Length Record	7
Theoretical Background	7
Exercise 2.1 — Encoding a Variable-Length Record	7
Exercise 2.2 — Handling NULL Fields	9
3 Lab 9-Part 3: Slotted Page for Variable-Length Records	10
4 Lab 9-Part 4: Free-Space Map (FSM)	11
5 Lab 9-Part 5: Buffer Replacement Policies	12
Theoretical Background	12
Exercise 5.1 — Implement <code>FIFOBuffer.access()</code>	12
Exercise 5.2 — Implement <code>LRUBuffer.access()</code>	13
Exercise 5.3 — Implement <code>RandomBuffer.access()</code>	13
Exercise 5.4 — Step-by-step Simulation (Same Trace, Same Capacity)	14

1 Lab 9. Part 1: Fixed-Length Record

Theoretical Background

A **fixed-length record** file stores all tuples with the same byte size. Each field is allocated its maximum length, so every record has an identical structure and predictable offset. This design simplifies addressing and I/O operations because the i -th record starts exactly at:

$$\text{offset}(i) = i \times \text{RECORD_SIZE}.$$

Example Schema

Consider an instructor record with the following schema:

```
(ID: varchar(5), name: varchar(20), dept_name: varchar(20), salary: numeric(8,2))
```

Instructor Record Structure

Assuming 1 byte per character and 8 bytes for the numeric field, the maximum record length is $5 + 20 + 20 + 8 = 53$ bytes.

In this lab, the schema is simplified to fixed-size ASCII fields and one integer salary field:

- **ID** – 5 bytes (ASCII)
- **Name** – 20 bytes (ASCII)
- **Department** – 20 bytes (ASCII)
- **Salary** – 4 bytes (little-endian integer, in cents)

Each record therefore occupies **49 bytes** in total.

Record Size and Page Capacity

Page Capacity Calculation

A single database page is 4KB (4096 bytes). The number of records that fit entirely within one page is:

$$\text{CAPACITY} = \left\lfloor \frac{\text{PAGE_SIZE}}{\text{RECORD_SIZE}} \right\rfloor = \left\lfloor \frac{4096}{49} \right\rfloor = 83.$$

Hence, each 4KB page can hold up to **83 complete records** without crossing page boundaries.

Code Implementation

```
1 PAGE_SIZE    = 4096    # 4KB page
2
3 # Structured dtype with fixed-length fields:
4 #   id: 5 bytes ASCII
5 #   name: 20 bytes ASCII
6 #   dept: 20 bytes ASCII
7 #   salary: 4-byte little-endian int (we store cents to avoid float issues)
8 dtype        = np.dtype([
9     ('id', 'S5'), ('name', 'S20'), ('dept', 'S20'), ('salary', '<i4')
10])
11 RECORD_SIZE = dtype.itemsize    # 5 + 20 + 20 + 4 = 49 bytes
12 CAPACITY    = PAGE_SIZE // RECORD_SIZE # number of records the page can hold
```

Listing 1: Definition of fixed-length record structure

The free-list method keeps a header pointer to the first free slot; each deleted slot stores the index of the next free one. On insertion, the head slot is reused first. This structure exactly matches the “free-list file” described in database textbooks (Figure 13.4).

Field Validation and Encoding Utilities

```
1 def _check_len(self, s: str, n: int, field: str)
2 def _encode(self, id5: str, name20: str, dept20: str, salary_float: float)
3 def _decode_row(self, rec)
```

Listing 2: Utility function definitions

Explanation of Utility Functions

_check_len(s, n, field) Ensures that a string value does not exceed the schema-defined maximum byte length. For example, the `id` field must be at most 5 bytes, while `name` and `dept` are limited to 20 bytes each. This keeps all records aligned to a fixed length and prevents overflow.

_encode(id5, name20, dept20, salary_float) Converts input values from Python objects into the binary form stored in the page. Strings are encoded as ASCII byte arrays, and the salary is scaled to integer cents and stored as a 4-byte integer. The function returns a tuple of encoded values ready for insertion into the NumPy structured array.

_decode_row(rec) Transforms a stored binary record back into a Python-friendly dictionary. It removes padding zeros from ASCII strings and divides the stored integer salary by 100 to recover the original floating-point value. This allows readable inspection of data when printing or debugging the page content.

Exercise 1.1 — Implement `insert()` (Dense Prefix)

Run: `python lab9_1_fixed_record_lab.py`

Exercise 1.1 — Implement `insert()` (dense-prefix append)

Goal: Append a new record at the end of the dense prefix $[0..count - 1]$ with no gaps.

```
1 def insert(self, id5: str, name20: str, dept20: str, salary_float: float):
2     """
3     Append a new record at the end of the dense prefix [0..count-1].
4     Assumes there are no holes in the current prefix.
5     """
6
7     # TODO:
8     # 1) If full (self.count >= CAPACITY), return None.
9     # 2) i = self.count
10    # 3) rid, rname, rdept, rcents = self._encode(id5, name20, dept20,
11        # salary_float)
12    # 4) Write fields into self.data at index i
13    # 5) self.count += 1
14    # 6) self.next_idx = max(self.next_idx, self.count)
15    # 7) return i
```

Listing 3: Function to implement: insert into dense prefix

Hints / Checkpoint

- Use `_encode(...)` to get fixed-size binary fields.
- After success, `count` increases by 1; `next_idx` stays $\geq count$.
- Return the index `i` where the new record was written.

Exercise 1.2 — Implement `delete_swap()`

Exercise 1.2 — Implement `delete_swap()` ($O(1)$, changes order)

Goal: Delete index i by moving the last record into position i .

```
1 def delete_swap(self, i: int) -> bool:
2     """
3     O(1) deletion by swapping the last record into position i.
4     This operation changes record order.
5     """
6
7     # TODO:
8     # 1) Validate  $0 \leq i < self.count$ ; else return False
9     # 2) last = self.count - 1
10    # 3) If  $i \neq last$ : self.data[i] = self.data[last]
11    # 4) self.count -= 1
12    # 5) self.next_idx = max(self.next_idx, self.count)
13    # 6) return True
```

Listing 4: Function to implement: $O(1)$ deletion using swap

Concept: Swap Deletion

- Complexity: **$O(1)$** — only overwrites one record.
- Changes logical order — record at the end moves into the deleted slot.
- Fast and suitable when record ordering is not important.

Exercise 1.3 — Implement `delete_compact()`

Exercise 1.3 — Implement `delete_compact()` (stable, $O(n)$)

Goal: Delete index i and preserve order by shifting subsequent records left.

```
1 def delete_compact(self, i: int) -> bool:
2     """
3     Stable deletion by shifting subsequent records left by one position.
4     Preserves order but requires more data movement.
5     """
6
7     # TODO:
8     # 1) Validate  $0 \leq i < self.count$ ; else return False
9     # 2) If  $i < self.count - 1$ :
10        self.data[i:self.count - 1] = self.data[i + 1:self.count]
11    # 3) self.count -= 1
12    # 4) self.next_idx = max(self.next_idx, self.count)
13    # 5) return True
```

Listing 5: Function to implement: stable deletion by compaction

Concept: Compact Deletion

- Complexity: **$O(n)$** — must shift all following records.
- Order-preserving, suitable for sequential scans.
- More expensive for large pages or frequent deletions.

Exercise 1.4 — Run and Observe (Practical Demo)

Exercise 1.4.1 — Insert Records (Dense Prefix)

Goal: Insert all instructor records and inspect the page state.

```
1 page.insert("10101", "Srinivasan", "Comp. Sci.", 65000.00)
2 page.insert("12121", "Wu", "Finance", 90000.00)
3 page.insert("15151", "Mozart", "Music", 40000.00)
4 page.insert("22222", "Einstein", "Physics", 95000.00)
5 page.insert("32343", "El Said", "History", 60000.00)
6 page.insert("33456", "Gold", "Physics", 87000.00)
7 page.insert("45565", "Katz", "Comp. Sci.", 75000.00)
8 page.insert("58583", "Califeri", "History", 62000.00)

9
10 print("stats:", page.stats())
11 page.dump_visible()
```

Checkpoint

Expect:

- `count = 8, next_idx = 8, free_head = -1`
- `free_chain = []` (no deleted slots)
- Visible indices 0–7 in insertion order

Exercise 1.4.2 — O(1) Deletion by Swap

Goal: Delete index 0 and observe reordering.

```
1 # page.delete_swap(0)
2 # print("after delete_swap(0), read(0):", page.read(0))
3 # print("stats:", page.stats())
4 # page.dump_visible()
```

Observe

The record at the last index overwrites slot 0. `count` decreases by 1; order changes. `next_idx` remains $\geq \text{count}$.

Exercise 1.4.3 — Stable Deletion (Compaction)

Goal: Delete index 0 again but preserve order.

```
1 # page.delete_compact(0)
2 # page.dump_visible()
3 # print("stats:", page.stats())
```

Observe

Records from indices [1..7] shift left; `count` decreases by 1, order preserved. Higher cost than swap deletion.

Exercise 1.4.4 — Free-list Deletion (Pointer-based)

Goal: Use a linked free-list for deletions without moving data.

```
1 # page.delete_to_freelist(1)
2 # page.delete_to_freelist(3)
3 # print("after delete_to_freelist(1,3):", page.stats())
4 # page.dump_visible()
```

Observe

Expect `free_head = 3, free_chain = [3, 1]`. No data movement — faster, but uses pointer links.

Exercise 1.4.5 — Free-list Insertion (Slot Reuse)

Goal: Reuse deleted slots via `insert_using_freelist()`.

```
1 # posA = page.insert_using_freelist("A0001", "Alice", "X", 11111.00)
2 # posB = page.insert_using_freelist("A0002", "Bob", "Y", 22222.00)
3 # posC = page.insert_using_freelist("A0003", "Carol", "Z", 33333.00)
4 # print("reused/append positions:", posA, posB, posC)
5
6 # page.dump_visible()
7 # print("final stats:", page.stats())
```

Observe

Slots are reused in **LIFO** order: first 3, then 1, then append to end. After reuse, `free_head = -1`, and `next_idx` advances only on append.

2 Lab 9-Part 2: Variable-Length Record

Theoretical Background

Unlike fixed-length records, variable-length records allow some fields (such as `name` or `dept`) to take different sizes in each record. A special header area at the beginning of each record stores metadata — including offsets, lengths, and a null bitmap — describing how to locate each field’s data within the payload area.

Record Structure Example

Schema:

```
(ID: varchar(5), Name: varchar(20), Dept: varchar(20), Salary: numeric(8,2))
```

Offsets in the header specify where each field begins and ends inside the record’s payload.

```
1 # Header format:
2 # - B: null bitmap (1 byte; bit0=id, bit1=name, bit2=dept, bit3=salary)
3 # - 6H: (off, len) pairs for id, name, dept (2 bytes each, unsigned short)
4 # - Q: salary (8 bytes unsigned integer; or use i32 if storing cents)
5
6 HDR_FMT = "<BHHHHHHQ"    # little-endian layout
7 HDR_SIZE = struct.calcsize(HDR_FMT)
8 print("HDR_SIZE", HDR_SIZE)
```

Listing 6: Header format and size calculation

Explanation of Header Fields

- **Null Bitmap (1 byte)** – each bit marks whether a field is NULL.
- **Offsets and Lengths (6 × 2 bytes)** – record the start position and length for each text field (ID, Name, Dept).
- **Salary (8 bytes)** – numeric value stored as an unsigned 64-bit integer (or 4-byte int for cents).
- **Total header size:** $1 + 12 + 8 = 21$ bytes.

The header serves as a miniature directory for the record, allowing the DBMS to locate each field efficiently, even when variable-length or NULL.

Exercise 2.1 — Encoding a Variable-Length Record

```
1 def encode_record(id_s, name_s, dept_s, salary, nulls=()):
2     """
3     Encodes a variable-length instructor record into binary form.
4     """
5     # Step 1. Build a 1-byte "null bitmap"
6     nb = 0
7     if 'id' in nulls:      nb |= (1 << 0)
8     if 'name' in nulls:   nb |= (1 << 1)
9     if 'dept' in nulls:   nb |= (1 << 2)
10    if 'salary' in nulls: nb |= (1 << 3)
11
12    # Step 2. Convert non-null fields into ASCII bytes
13    b_id   = b"" if (nb & (1 << 0)) else _b(id_s)
14    b_name = b"" if (nb & (1 << 1)) else _b(name_s)
15    b_dept = b"" if (nb & (1 << 2)) else _b(dept_s)
16
17    # Step 3. Concatenate all variable-length text fields
18    payload = b_id + b_name + b_dept
```

```

19
20     # Step 4. Compute field offsets relative to start-of-record
21     # TODO:
22     # off_id    = ???
23     # off_name  = ???
24     # off_dept  = ???
25
26     ## TODO: Print offsets, lengths, and values for debugging
27     print("----Record/Layout(Header|Summary)----")
28     print("ID->offset=? , length=? , value=?")
29     print("Name->offset=? , length=? , value=?")
30     print("Dept->offset=? , length=? , value=?")
31
32     # Step 5. Numeric salary (0 if NULL)
33     sal = 0 if (nb & (1 << 3)) else int(salary)
34
35     # Step 6. Pack the header
36     header = struct.pack(
37         HDR_FMT,
38         nb,
39         off_id, len(b_id),
40         off_name, len(b_name),
41         off_dept, len(b_dept),
42         sal
43     )
44
45     # Step 7. Return concatenation of header and payload
46     return header + payload

```

Listing 7: Function to implement: encode a variable-length record

Hints for Students

- Step 4 computes offsets relative to the record start. The first payload byte begins at `HDR_SIZE`.
- The `print(...)` lines are used to check correctness: verify offsets and lengths match each field's byte size.
- Step 6 packs all metadata into the fixed-size header.
- The final record is `header + payload`.

Exercise 2.2 — Handling NULL Fields

Test NULL Field Encoding

Goal: Encode records with some fields set to NULL using the `nulls` argument.

```
1 # Example 1: No NULLs
2 r1 = encode_record("10101", "Srinivasan", "Comp. Sci.", 6500000)
3
4 # Example 2: Only salary is NULL
5 r2 = encode_record("10101", "Srinivasan", "Comp. Sci.", 0, nulls={"salary"})
6
7 # Example 3: Multiple NULLs (name and dept)
8 r3 = encode_record("10101", "", "", 7000000, nulls={"name", "dept"})
```

Checkpoint

- Observe how the null bitmap (`nb`) changes:
 - All non-NUL then 0000 (0)
 - Salary NULL then 1000 (8)
 - Name + Dept NULL then 0110 (6)
- Note that payload size shrinks when more fields are NULL.

Try Modifying a Field to NULL

Goal: Modify an existing record to make one or more fields NULL and observe the change in offsets.

```
1 # Original record (no NULL)
2 r_orig = encode_record("12121", "Wu", "Finance", 9000000)
3
4 # Modified record (dept becomes NULL)
5 r_null = encode_record("12121", "Wu", "", 9000000, nulls={"dept"})
```

Observation

Compare the printed offsets of `r_orig` and `r_null`. The missing field shortens the payload and shifts subsequent offsets.

Concept Summary

- Variable-length records combine a fixed-size header and a flexible payload.
- Offsets and lengths allow direct access to each field's data.
- A 1-byte null bitmap efficiently marks missing attributes.
- This structure is the foundation for the slotted-page layout used in modern DBMSs.

3 Lab 9-Part 3: Slotted Page for Variable-Length Records

Theoretical Background

A **slotted page** manages variable-length records within a fixed-size page. Two regions grow toward each other:

- The slot directory (array of record pointers) grows upward.
- The data region (actual records) grows downward.

Free space lies between them. Each slot stores (offset, length) for its record.

Core Code

```
1 def insert_record(self, rec_bytes):
2     size = len(rec_bytes)
3     free_start = (len(self.slots)+1) * 8
4     if size + 8 > self.free_end - free_start: return None
5     self.free_end -= size
6     offset = self.free_end
7     self.data[offset:offset+size] = rec_bytes
8     self.slots.append((offset, size))
```

Listing 8: Insert into slotted page

Explanation

Slots provide logical record identifiers. Deleting a record invalidates its slot rather than moving data. Free space fragmentation can be reclaimed by compaction when necessary.

Demo

Run: `python lab9_3_slotted_page_lab.py`

Expected: records inserted with varying lengths, slot table updated dynamically.

4 Lab 9-Part 4: Free-Space Map (FSM)

Theoretical Background

A Free-Space Map (FSM) tracks the amount of unused space per data page. A two-level FSM reduces search cost:

- Level-0: exact free-space buckets per page,
- Level-1: max free fraction per group of pages.

Core Code

```
1 def find_page(self, min_free_frac=0.0):
2     need = self.bucketize(min_free_frac)
3     for i, mx in enumerate(self.l1):
4         if mx >= need:
5             gidx = i; break
6     s, e = gidx*self.G, min(len(self.l0), gidx*self.G+self.G)
7     for pid in range(s, e):
8         if self.l0[pid] >= need: return pid
```

Listing 9: Two-level FSM search algorithm

Explanation

By summarizing groups, FSM reduces the expected scan time from $O(N)$ to $O(\sqrt{N})$. This hierarchical structure is conceptually similar to a multi-level index.

Demo

Run: `python lab9_4_fsm_lab.py`

Observe how FSM efficiently locates pages with sufficient free space.

5 Lab 9-Part 5: Buffer Replacement Policies

Theoretical Background

A **database buffer pool** is a fixed-size memory area used to cache recently accessed disk pages. When the system needs to access a page, it first checks whether the page is already in the buffer:

- **HIT**: the page is found in memory, so no disk I/O is needed.
- **FAULT**: the page is not in memory and must be loaded from disk.

Because memory capacity is limited, some page must be **evicted** when the buffer is full and a new page arrives. A **replacement policy** determines which page to remove.

Page Replacement Concept

Each buffer access produces either a HIT or a FAULT. If a FAULT occurs and the buffer is full, one page is selected for eviction:

```
disk then buffer (if full then evict one page then insert new page)
```

Common replacement policies:

- **FIFO (First-In-First-Out)** — evict the page that has been in memory the longest.
- **LRU (Least Recently Used)** — evict the page that has not been accessed for the longest time.
- **Random** — evict a random page (simple baseline for comparison).

Exercise 5.1 — Implement FIFOBuffer.access()

```
1  from collections import deque
2
3  class FIFOBuffer:
4      """
5          FIFO (First-In-First-Out) page replacement.
6          - Keep pages in a queue (left = oldest, right = newest).
7          - On FAULT when full: evict the oldest page (front of queue).
8          - On HIT: order does not change.
9      """
10     def __init__(self, capacity: int):
11         self.capacity = capacity
12         self.q = deque()           # queue: leftmost is oldest
13         self.in_mem = set()       # O(1) membership
14         self.last_evicted = None # record last evicted page for printing
15
16     def access(self, page_id: int) -> bool:
17         """
18             Return True if FAULT, False if HIT.
19             Update self.last_evicted when eviction happens.
20         """
21         # TODO:
22         # 1) If page_id in self.in_mem:    # HIT
23
24         #
25         # 2) Else (FAULT) and self.capacity > 0:
26
27         pass
```

Listing 10: TODO: FIFO (queue-based) replacement

Hints

- Evict from the left of the queue (`popleft()`).
- Insert new pages at the right (`append()`).
- Always keep `in_mem` in sync with the queue.

Exercise 5.2 — Implement `LRUBuffer.access()`

```
1  from collections import OrderedDict
2
3  class LRUBuffer:
4      """
5          LRU (Least Recently Used) page replacement.
6          Use an OrderedDict to maintain recency.
7          On HIT: move page to end (most recent).
8          On FAULT when full: evict from front (least recent).
9      """
10     def __init__(self, capacity: int):
11         self.capacity = capacity
12         self.od = OrderedDict()    # keys in recency order: left=LRU, right=MRU
13         self.last_evicted = None
14
15     def access(self, page_id: int) -> bool:
16         """
17             Return True if FAULT, False if HIT.
18             Update self.last_evicted when eviction happens.
19         """
20         # TODO:
21         # 1) If page_id in self.od:           # HIT
22
23         #
24         # 2) Else (FAULT) and self.capacity > 0:
25
26         #
27         # 3) return True
28         pass
```

Listing 11: TODO: LRU (recency-based) replacement

Hints

- `move_to_end(page_id)` marks a page as most recently used.
- `popitem(last=False)` evicts the least recently used page.

Exercise 5.3 — Implement `RandomBuffer.access()`

```
1  import random
2
3  class RandomBuffer:
4      """
5          Random page replacement (baseline).
6          On FAULT when full: evict one random page.
7          Does not use recency or frequency.
8      """
9      def __init__(self, capacity: int):
10         self.capacity = capacity
11         self.pages = set()           # pages in memory (unordered)
12         self.last_evicted = None
13
14     def access(self, page_id: int) -> bool:
15         """
```

```

16     Return True if FAULT, False if HIT.
17     Update self.last_evicted when eviction happens.
18     """
19     # TODO:
20     # 1) If page_id in self.pages:    # HIT
21
22     # 2) Else (FAULT) and self.capacity > 0:
23
24     pass

```

Listing 12: TODO: Random (baseline) replacement

Hints

- Convert the set to a list for `random.choice`.
- Keep `last_evicted` as `None` when no eviction occurs.

Exercise 5.4 — Step-by-step Simulation (Same Trace, Same Capacity)

Goal

Run all three policies on the *same* short access sequence with the *same* capacity. Print each step: HIT/FAULT, evicted page, and the buffer state. This makes the behavioral differences obvious.

```

1 def buffer_order(buf):
2     """Pretty-print buffer state by policy type."""
3     if hasattr(buf, "q"):        # FIFO: oldest -> newest
4         return list(buf.q)
5     if hasattr(buf, "od"):       # LRU:   LRU -> MRU
6         return list(buf.od.keys())
7     if hasattr(buf, "pages"):   # Random: unordered set (sorted for stable view)
8         return sorted(buf.pages)
9     return []
10
11 def run_verbose(trace, capacity, policy_cls, policy_name):
12     """
13     Print per-step:
14     accessed page
15     HIT/FAULT
16     evicted page (if any)
17     buffer state after the access
18     """
19     buf = policy_cls(capacity)
20     print(f"\n==>{policy_name}(capacity={capacity})==")
21     print(f"{'Step':>4}{'Access':>6}{'Result':>6}{'Evicted':>8}{'State':>65}")
22     print("-" * 65)
23     faults = 0
24     for step, p in enumerate(trace, start=1):
25         result = buf.access(p)
26         evicted = getattr(buf, "last_evicted", None)
27         if result:
28             faults += 1
29         print(f"{step:4d}{p:6d}{('FAULT' if result else 'HIT'):>6}{str(evicted):>8}{buffer_order(buf)}")
30     print(f"Total faults: {faults}/{len(trace)} (Fault rate = {faults/len(trace):.3f})")
31
32 # Demo trace and capacity (short and readable)
33 trace = [1, 2, 3, 1, 4, 5, 2, 1, 2, 3, 4, 5]

```

```
35 capacity = 3
36
37 # Run all three policies
38 run_verbose(trace, capacity, FIFOBuffer, "FIFO")
39 run_verbose(trace, capacity, LRUBuffer, "LRU")
40 run_verbose(trace, capacity, RandomBuffer, "Random")
```

Listing 13: Step-by-step runner with a short trace

What to Observe

- **FIFO**: queue order is oldest to newest. Notice how old but still-useful pages can be evicted.
- **LRU**: recently used pages are retained; least recent is evicted.
- **Random**: victims change unpredictably; fault rate fluctuates.