
DTS207TC Database Development and Design

Lecture 6
UML for CW1Q4

Chap 10 Big Data*

Di Zhang, Autumn 2025

*Page titles with * will not be assessed*

Outline

- UML
- Big Data
- Hadoop

- (a) Design UML schema for this application and draw UML class diagram for the schema. Specify key attributes of each entity type and structural constraints on each relationship type. Note any unspecified requirements, and make appropriate assumptions to make the specification complete. (10 marks)
- (b) Translate the UML class diagram created in the above question (a) into a relational schema. In the relational schema, only use relation name and attribute, e.g., student (name, number, SSN,..etc.) and specify foreign key as fk. (10 marks)

This part belongs to object-oriented programming (C++, Python, etc.)

What is UML?

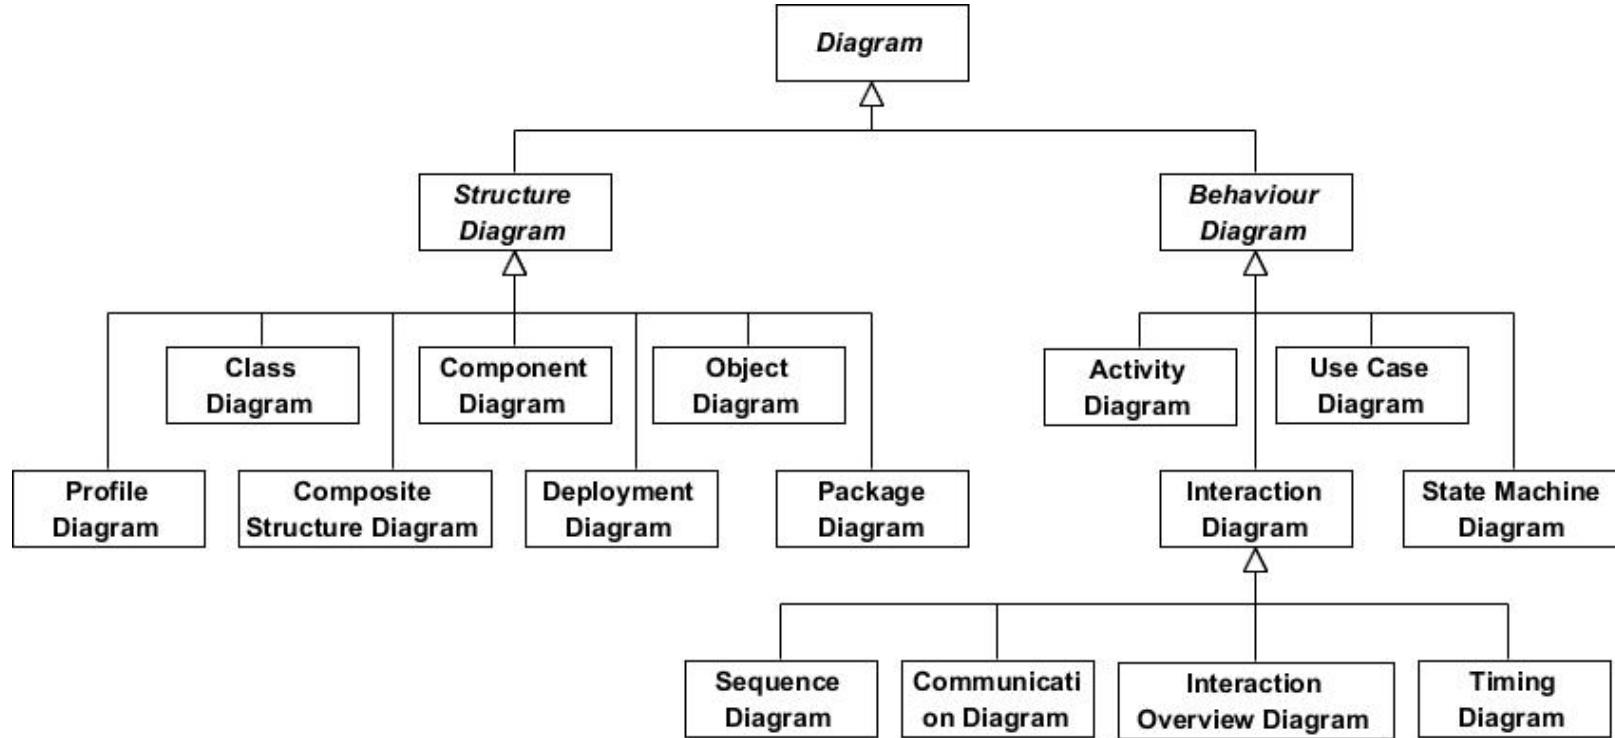
- Definition:
 - A standardized, powerful graphical modeling language.
 - Used to visualize, specify, construct, and document software-intensive systems.
- Core Purpose:
 - Communication: Provides a unified bridge for communication between development teams, clients, and architects.
 - Blueprint: Creates a blueprint for the system before coding, reducing complexity in understanding and design.
 - Documentation: Generates documentation crucial for understanding and maintaining the system.
- It is not a programming language, but a tool for designing software models.



Overview of UML Diagrams

- UML diagrams are primarily divided into Structure Diagrams and Behavior Diagrams.
- Structure Diagrams (Static Model) - Describe the static structure of a system.
 - **Class Diagram** - Core diagram showing classes, interfaces, and their static relationships.
 - Object Diagram - A snapshot of object instances and their relationships at a point in time.
 - Component Diagram - Describes the physical components and their dependencies.
 - Deployment Diagram - Shows how software is deployed on hardware.
 - Composite Structure Diagram - Describes the internal structure of a class.
 - Package Diagram - Used to group elements.
- Behavior Diagrams (Dynamic Model) - Describe the dynamic behavior of a system.
 - Use Case Diagram - Describes system functionality from a user's perspective.
 - Activity Diagram - Models workflows or business processes.
 - State Machine Diagram - Describes the state changes of an object during its life cycle.
 - Sequence Diagram - Emphasizes the time order of messages in an interaction.
 - Communication Diagram - Emphasizes the structural organization of objects in an interaction.

Overview of UML Diagrams



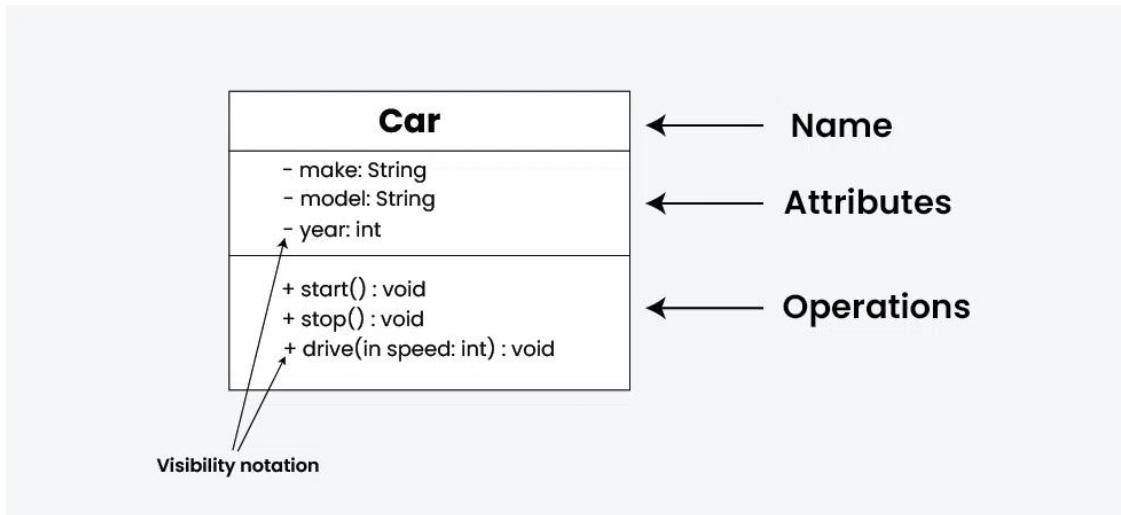
Class Diagrams - The Heart of UML

- What is a Class Diagram?
 - A diagram that describes the classes, interfaces, and their static relationships within a system.
 - The most commonly used and fundamental diagram for object-oriented system modeling.
- Main Purpose of Class Diagrams:
 - Provides a unified view of the system design for the development team.
 - Details the logical structure of the system.
 - Serves as the basis for Forward Engineering (generating code skeletons from class diagrams) and Reverse Engineering (generating class diagrams from code).
- Core Components: Classes, Attributes, Methods, Relationships.

The Basic Structure of a Class



- A class is represented in UML as a rectangle divided into three compartments:
 - ClassName
 - attributes
 - operations



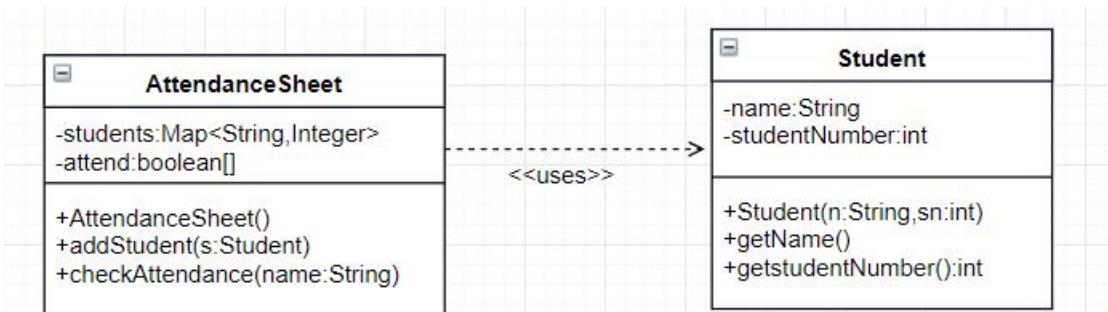
Class Relationships

- 1. Association
- Description: A structural relationship between classes, indicating a long-lasting, "knows-a" connection between objects.
- Notation: A solid line.
- Navigability: Arrows indicate the direction that can be navigated. A bidirectional association may have no arrows or two arrows.
- Role Names: Can be added at each end of the association to describe the role the other class plays.
- Multiplicity: Indicates how many objects of one class relate to another class.
- 1 (exactly one), 0..1 (zero or one), * or 0..* (many), 1..* (one or more), n (a specific number)



Class Relationships

- 2. Dependency
- Description: A "using" relationship, temporary and transient. A change in one class may affect the dependent class.
- Notation: A dashed line with an open arrow ----->
- Common Scenarios: Method parameters, local variables, static method calls.



Example: Association vs. Dependency

```
class Person:  
    """  
  
        Association Example: Person HAS-A Car.  
        This is a structural, long-term relationship.  
    """  
  
    def __init__(self, car: Car):  
        # Association: Person "has" a Car.  
        # 'my_car' is an instance variable (field).  
        # The relationship persists as long as the Person object exists.  
        self.my_car = car  
  
  
    def drive_my_car(self):  
        """Uses the car that the person owns."""  
        self.my_car.start()
```

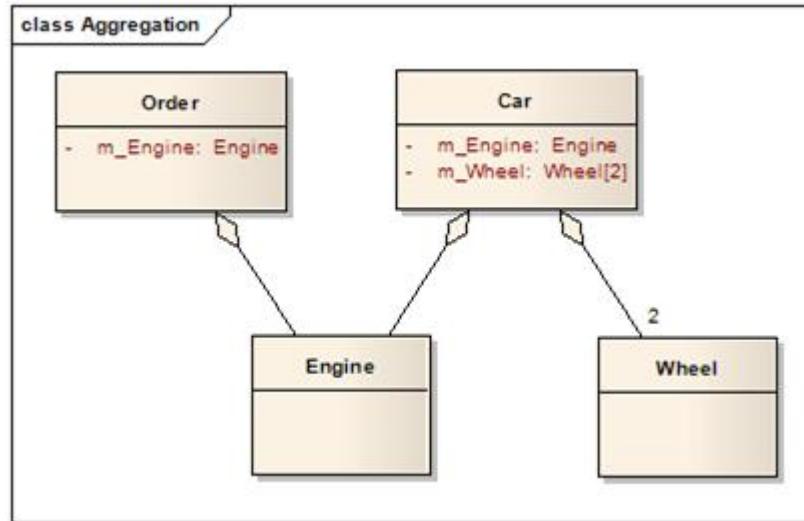
Example: Association vs. Dependency

```
def drive_rental_car(self, rental_car: Car):
    """
    Dependency: Person "uses" a rental car temporarily.
    'rental_car' is a method parameter, a temporary input.
    """
    rental_car.start()
    # The relationship with this specific 'rental_car' object ends when the method finishes.

def go_to_test_drive(self):
    """
    Another Dependency example: creating a local variable inside a method.
    """
    test_car = Car() # 'test_car' is a local variable
    test_car.start()
    # When the method ends, the 'test_car' reference goes out of scope.
```

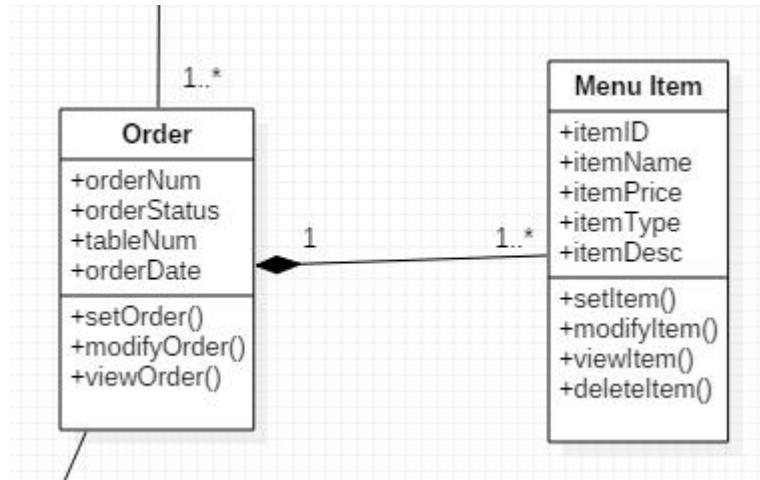
Class Relationships

- 3. Aggregation
- Description: A special type of Association representing a "whole-part" relationship where the part can exist independently of the whole. It's a weak "has-a" relationship.
- Notation: A solid line with a hollow diamond, diamond connected to the whole.
- Whole ◊ —— Part



Class Relationships

- 4. Composition
- Description: A stronger form of Aggregation, representing a strict "whole-part" relationship where the part's lifecycle is dependent on the whole. It's a strong "has-a" relationship.
- Notation: A solid line with a filled diamond, diamond connected to the whole.
- Whole ◆ — Part



Example: Aggregation

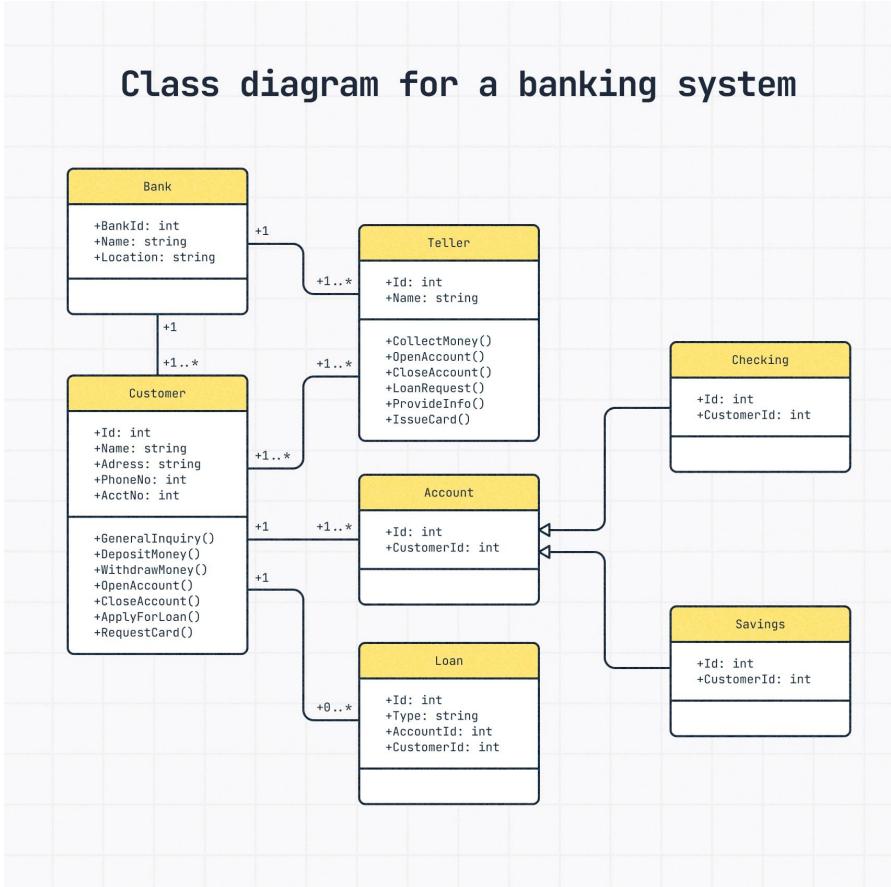
```
class Department:  
    """  
        Aggregation Example: Department HAS-A Teacher, but weakly.  
        The Teacher can exist independently of the Department.  
    """  
  
    def __init__(self, name: str):  
        self.name = name  
        # Aggregation: Department has a list of Teachers  
        # The teachers are passed from outside and can exist without the department  
        self.teachers = [] # This is a weak "has-a" relationship
```

Example: Composition

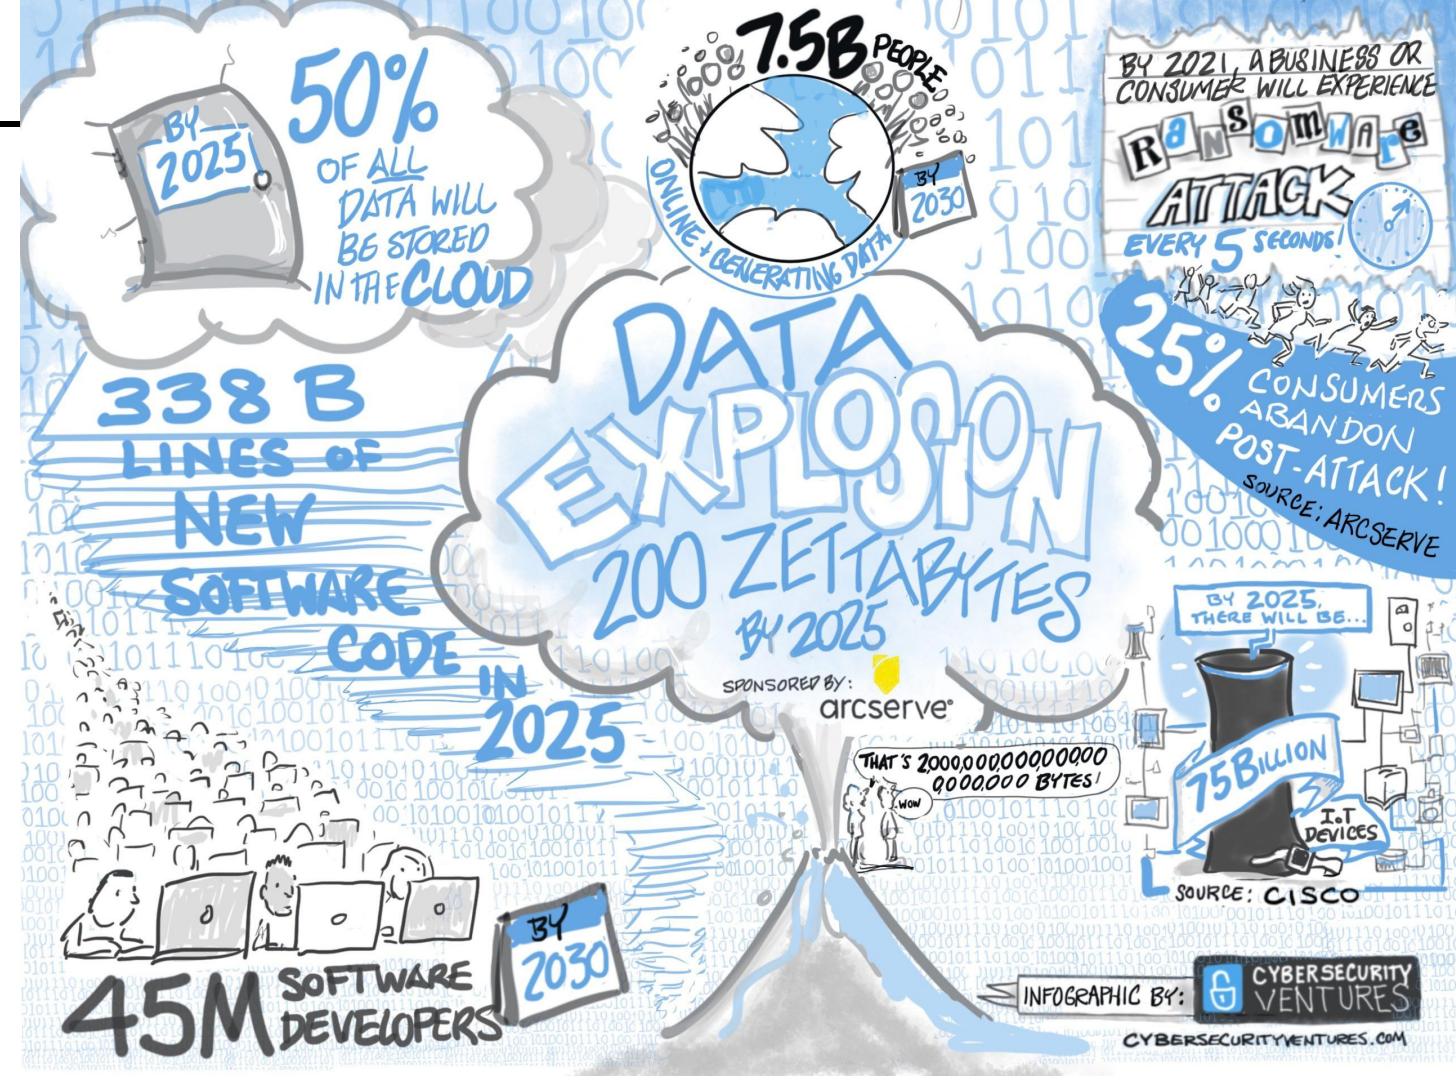
```
class Car:  
    """  
        Composition Example: Car HAS-A Engine and Wheels, STRONGLY.  
        The Engine and Wheels are created when the Car is created and cannot exist without it.  
    """  
  
    def __init__(self, model: str):  
        self.model = model  
  
        # Composition: Car creates its own Engine and Wheels  
        # These parts are created WHEN the car is created and have no independent existence  
        self.engine = Engine("V6") # Strong "has-a" relationship  
        self.wheels = [  
            Wheel("Front Left"), # Strong "has-a" relationship  
            Wheel("Front Right"), # Strong "has-a" relationship  
            Wheel("Rear Left"), # Strong "has-a" relationship  
            Wheel("Rear Right") # Strong "has-a" relationship  
        ]  
        print(f"Car {self.model} created with engine and 4 wheels")
```

Example

Class diagram for a banking system



- Any drawing tool is acceptable. Such as:
 - <https://plantuml.com/>
 - https://www.lucidchart.com/pages/examples/uml_diagram_tool
 - <https://www.visual-paradigm.com/solution/freeumltool/>
 - or whatever tool you choose



BIG DATA

Big Data is data that is too large, complex and dynamic for any conventional data tools to capture, store, manage and analyze.

The right use of *Big Data* allows analysts to spot trends and gives niche insights that help create value and innovation much faster than conventional methods.



57.6% OF ORGANIZATIONS SURVEYED SAY THAT BIG DATA IS A CHALLENGE



72.7% CONSIDER DRIVING OPERATIONAL EFFICIENCIES TO BE THE BIGGEST BENEFIT OF A BIG DATA STRATEGY

50% SAY THAT BIG DATA HELPS IN BETTER MEETING CONSUMER DEMAND AND FACILITATING GROWTH

The “three V’s”, i.e the Volume, Variety and Velocity of the data coming in is what creates the challenge.

VOLUME

Region	Amount of Big Data stored (in petabytes)
North America	>3,500
Europe	>2,000
China	>250
Japan	>400
Middle East	>200
India	>50
Latin America	>50

VARIETY

PEOPLE TO PEOPLE

NETIZENS, VIRTUAL COMMUNITIES, SOCIAL NETWORKS, WEB LOGS...

PEOPLE TO MACHINE

ARCHIVES, MEDICAL DEVICES, DIGITAL TV, E-COMMERCE, SMART CARDS, BANK CARDS, COMPUTERS, MOBILES...

MACHINE TO MACHINE

SENSORS, GPS DEVICES, BAR CODE SCANNERS, SURVEILLANCE CAMERAS, SCIENTIFIC RESEARCH...

VELOCITY

2.9 MILLION

EMAILS SENT EVERY SECOND

20 HOURS

OF VIDEO UPLOADED EVERY MIN

50 MILLION

TWEETS PER DAY

CASE STUDY - Healthcare

\$300 billion is the potential annual value to Healthcare

\$165B
CLINICAL

\$9B
PUBLIC HEALTH

\$108B
R&D

\$5B
BUSINESS MODEL

\$47B
ACCOUNTS

TRANSPARENCY IN CLINICAL DATA AND CLINICAL DECISION SUPPORT

AGGREGATION OF PATIENT RECORDS, ONLINE PLATFORMS AND COMMUNITIES

PUBLIC HEALTH SURVEILLANCE AND RESPONSE SYSTEMS

ADVANCED FRAUD DETECTION; PERFORMANCE BASED DRUG PRICING

RESEARCH AND DEVELOPMENT; PERSONALIZED MEDICINE; CLINICAL TRIAL DESIGN

VALUE

Industry	PRODUCTIVITY INCREASE (%)	SALES INCREASE (\$)
Retail	49%	\$9.6B
Consulting	39%	\$5.0B
Air Transportation	21%	\$4.3B
Construction	20%	\$4.2B
Food Products	20%	\$3.4B
Steel	20%	\$3.4B
Automobile	19%	\$2B
Industrial Instruments	18%	\$1.2B
Publishing	18%	\$0.8B
Telecommunications	17%	\$0.4B

40% PROJECTED GROWTH IN GLOBAL DATA CREATED PER YEAR

5% PROJECTED GROWTH IN GLOBAL IT SPENDING PER YEAR

The estimated size of the digital universe in 2011 was 1.8 zettabytes. It is predicted that between 2009 and 2020, this will grow 44 fold to 35 zettabytes per year. A well defined data management strategy is essential to successfully utilize Big Data.

Sources: ① Realizing the Rewards of Big Data - Wipro Report ② Big Data: The Next Frontier for Innovation, Competition and Productivity - McKinsey Global Institute Report ③ Stream, Redshift Group ④ Measuring the Business Impacts of Effective Data - study by University of Texas, Austin ⑤ US Department of Labour: NYSE-WRT | OVER 130,000 EMPLOYEES | 154 COUNTRIES | CONSULTING | SYSTEM INTEGRATION | OUTSOURCING

DO BUSINESS BETTER

WIPRO
Applying Thought

Visualization: Senthil

University
大學

6 Steps in **CRISP-DM**

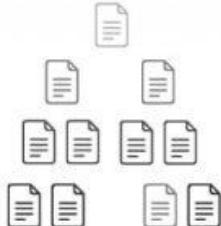
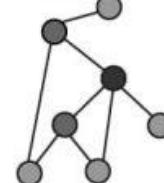
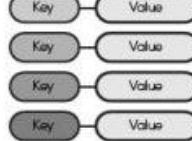
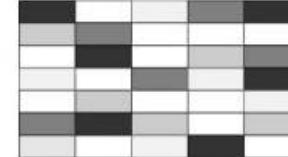
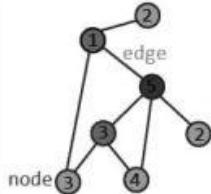
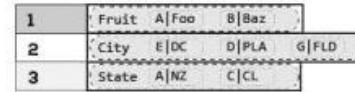
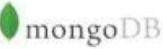
The Standard
Data Mining
Process



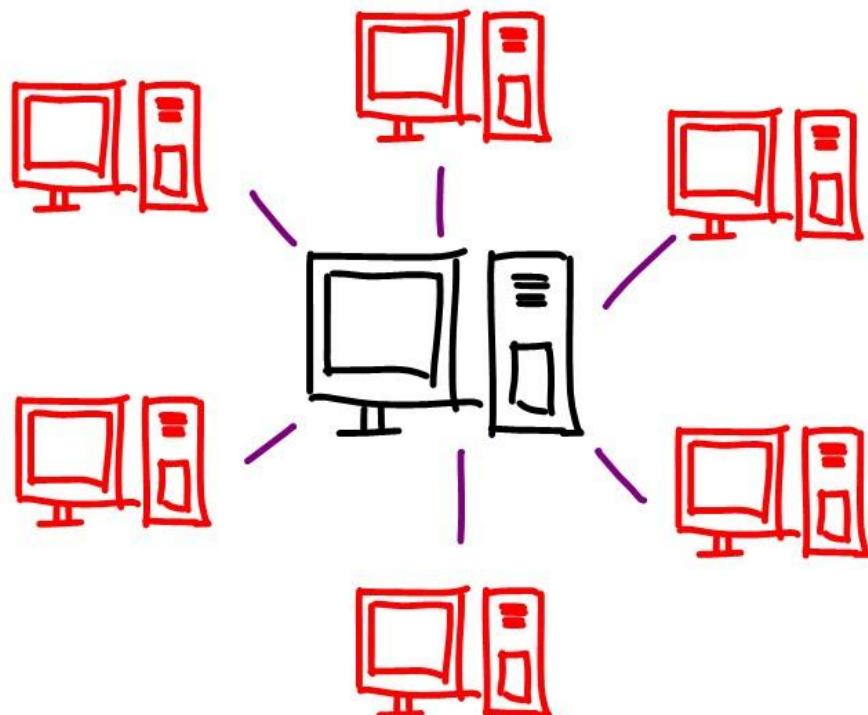
Challenge and Response

- Volume
 - Distributed Computing
- Velocity
 - ACID->BASE: NoSQL
- Variety
 - Unstructured Text

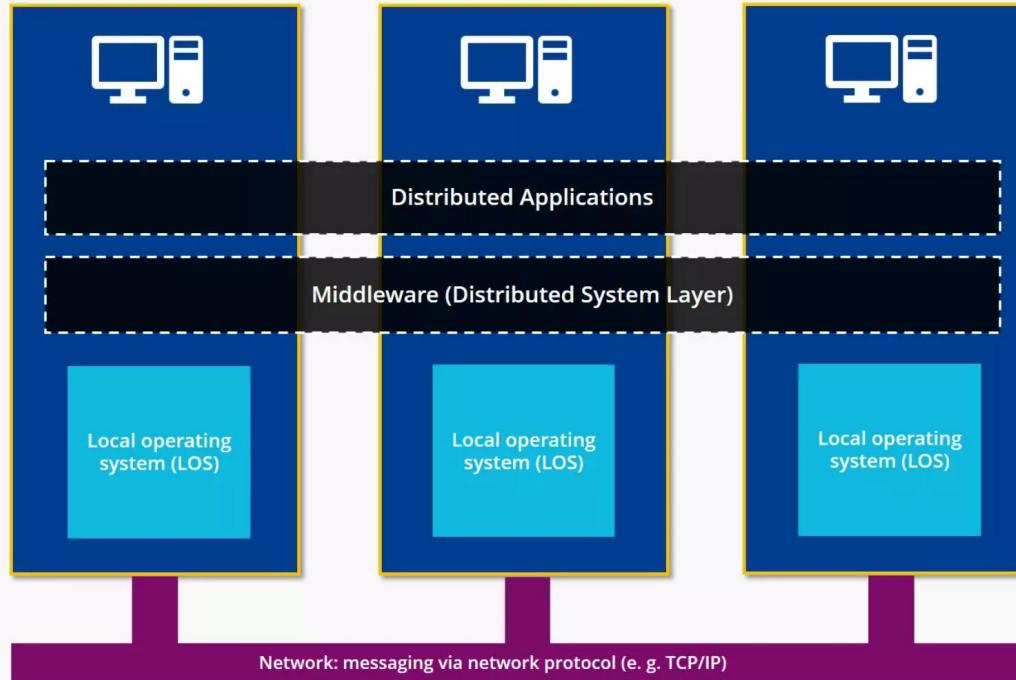
NoSQL DATABASE TYPES

Document	Graph	Key-Value	Wide-Column
			
<pre>{ "user":{ "id":"143", "name":"improgrammer", "city":"New York" } }</pre>			
  	 	   	  

Distributed Computing



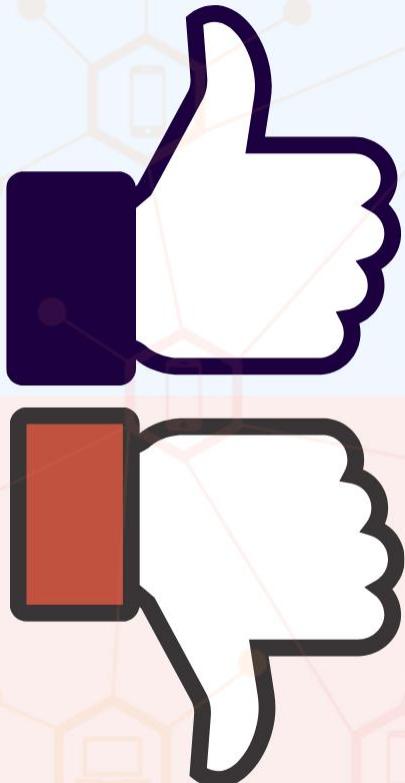
Distributed Computing



IONOS

Distributed applications can solve problems across devices in a computer network. When used in conjunction with middleware, they can optimize operational interactions with locally accessible hardware and software.

Advantages



🔒 Improved Security

💡 Flexibility and Scalability

🚀 Enhanced Performance

🔗 Increased Reliability

💾 Expanded Storage

\$ Cost effectiveness

⚡ Low Latency

Distributed Computing

💰 Maintenance Fees

💔 Component Failure

🌐 Bandwidth Restrictions

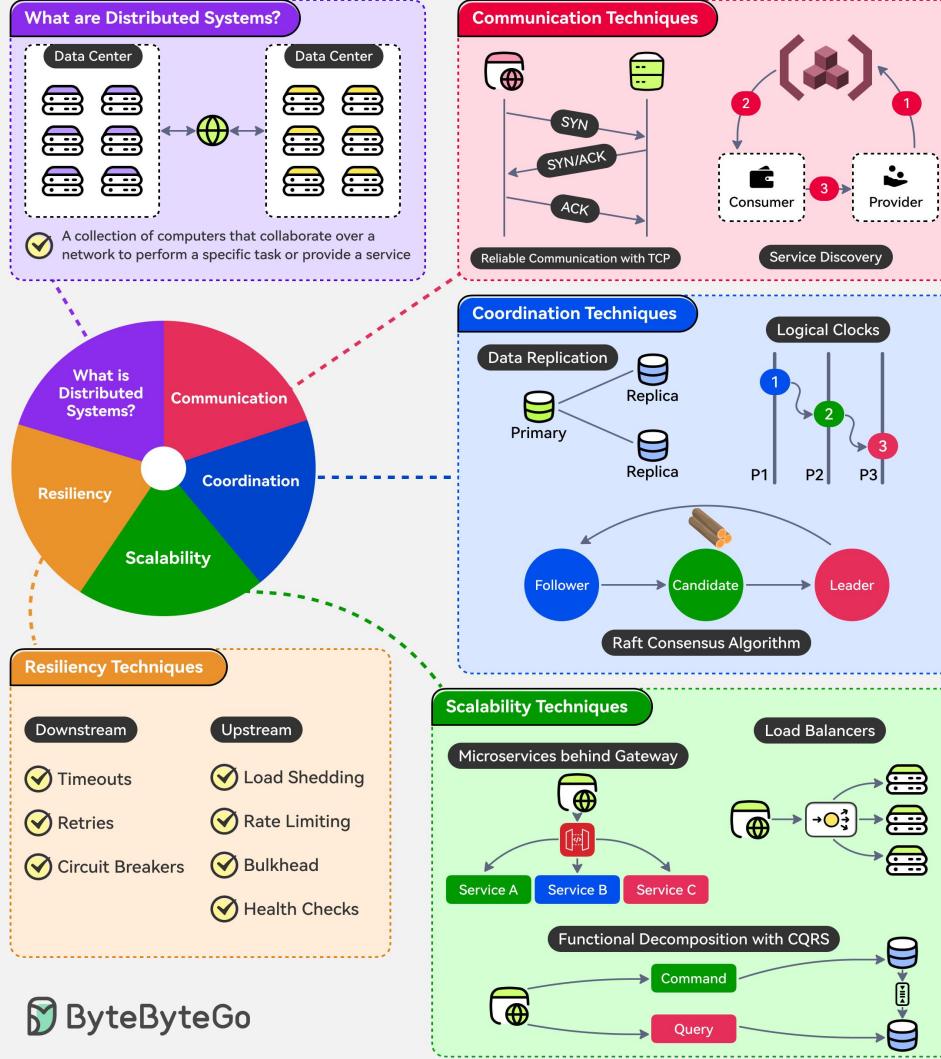
🚧 Development Obstacles

🐢 Slow Network Transfers

❓ Issues with Standardization

✳️ Systems Complexity

Disadvantages



Distributed System

- A distributed system is a collection of independent computers that appear to users as a single coherent system
- Focuses on providing reliability, availability, & fault tolerance by distributing data & processes across multiple nodes
- To ensure system resilience, fault tolerance, & scalability by distributing services or data
- Involves systems such as distributed databases, message queues, or microservices

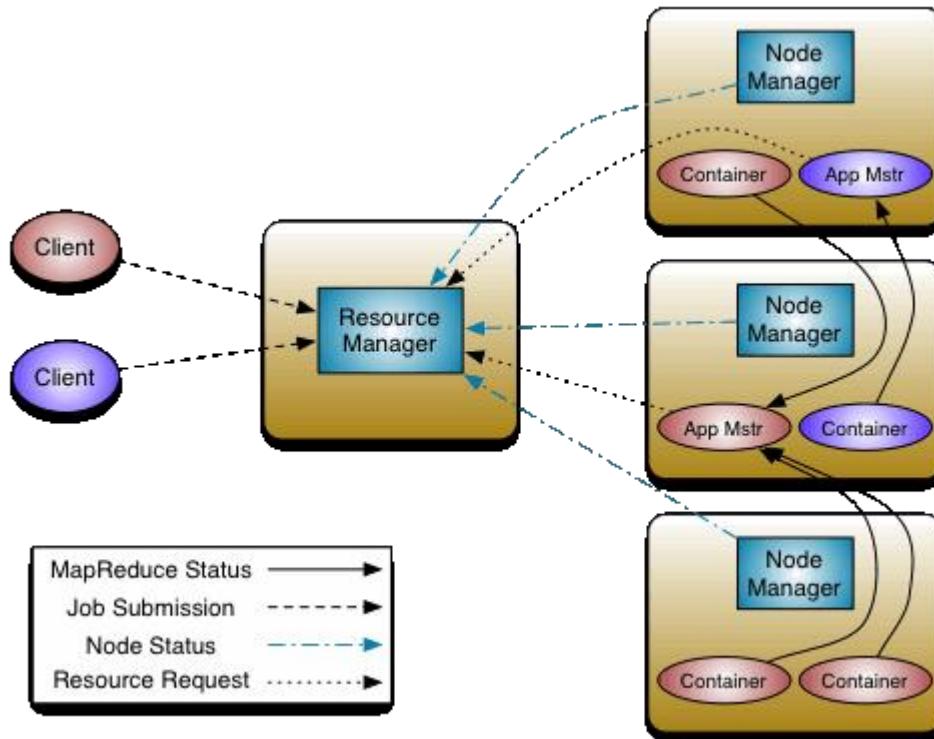


Distributed Computing

- Distributed computing refers to using multiple computers to perform computational tasks collaboratively.
- Focuses on splitting large computational tasks into smaller chunks that can be executed concurrently.
- To solve complex computational problems by breaking them down and distributing tasks across multiple nodes.
- Involves dividing a large computational workload into smaller tasks

Aspect	Parallel Computing	Distributed Computing
Primary Goal	Increase speed, reduce execution time for a single, complex task.	Resource sharing, scalability, fault tolerance, solving large-scale problems.
System Architecture	Tightly-coupled. Processors connected via high-speed buses, shared memory, or specialized networks (InfiniBand).	Loosely-coupled. Computers (nodes) connected via standard networks (Ethernet, Internet).
Memory Model	Often Shared Memory. All processors access a common memory address space.	Often Distributed Memory. Each node has its own private memory; data is exchanged via message passing (e.g., MPI).
Communication & Coordination	Very low latency, very high bandwidth. Processes/threads require tight synchronization.	Higher latency, lower bandwidth. Nodes collaborate through asynchronous message passing.
User Perspective	Feels like using a single, powerful supercomputer.	Feels like using a single, coherent system (even though it's composed of many machines).
Scaling Method	Vertical Scaling (Scale-up). Adding more CPUs/cores within a single node. Limited by physical constraints.	Horizontal Scaling (Scale-out). Adding more commodity computer nodes. Theoretically unlimited.
Fault Tolerance	Typically weak. The failure of a single processor can cause the entire computation to fail.	Inherently strong. Designed with node failure in mind. Tasks can be restarted on other nodes if one fails.
Typical Applications	Scientific computing, 3D rendering, weather simulation, genomic sequence analysis.	Large web services (Google, Amazon), Big Data processing (Hadoop/Spark), blockchain, cloud platforms.
Programming Models/Tools	OpenMP (shared memory), MPI (distributed memory, for tight clusters), CUDA (GPU).	Hadoop, Spark, Akka, gRPC, and various cloud APIs (AWS, Azure).

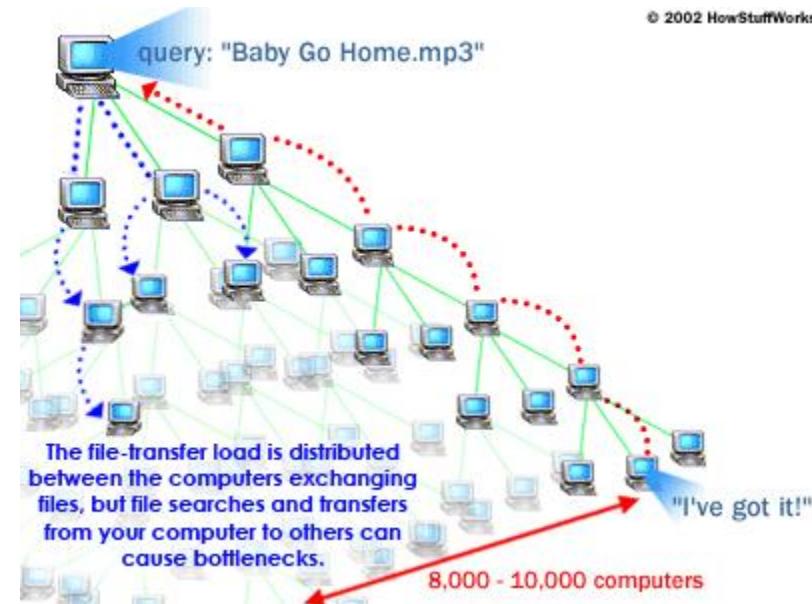
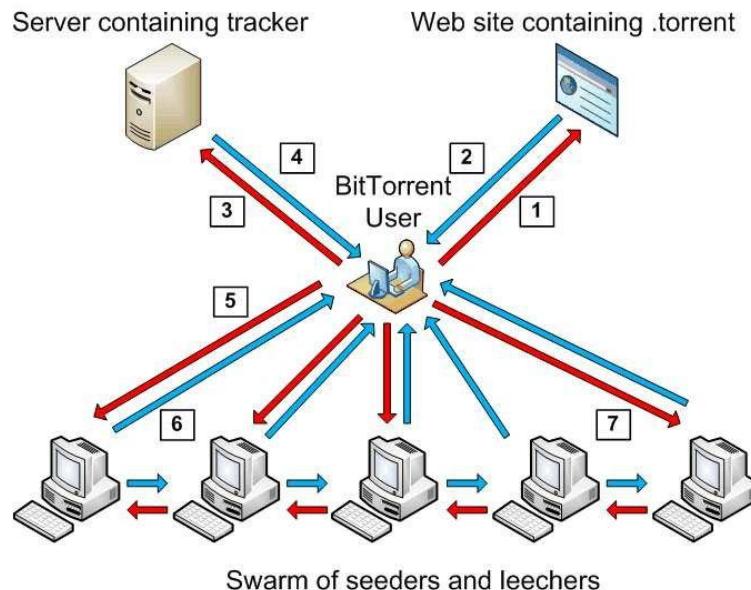
Example: Hadoop



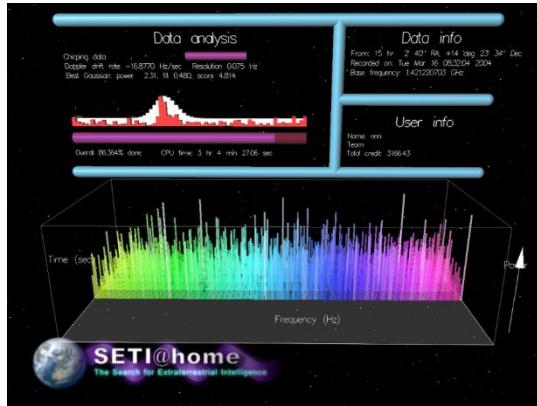
Why Hadoop belongs to distributed Computing, but MPI not?

- 1. Primary Goal: Scale vs. Speed
 - Extreme Speed vs. the Ability to Process Massive Data
- 2. System Architecture: Fault Tolerance vs. Performance (the most critical difference)
 - Parallel Computing (MPI): Tightly coupled. After tasks are distributed to multiple nodes, frequent communication and synchronization between nodes are required.
 - Hadoop: Loosely coupled. Its core design philosophies are "mobile computing is cheaper than moving data" and "hardware failures are the norm, not the exception."
- 3. Communication Model: Message Passing vs. Data Sharing
 - Parallel Computing (MPI): Frequent, low-latency, bidirectional communication between processes is achieved through a message passing interface.
 - Hadoop MapReduce: There is almost no communication between processes. This "share nothing" architecture is a typical distributed system design.
- 4. Problem Domain: Single Complex Task vs. Decomposable Batch Tasks
 - Parallel Computing: Excels at solving a single, complex problem that requires a high degree of coordination (such as simulating an entire weather system). Hadoop MapReduce excels at batch processing problems that can be clearly decomposed into a large number of independent subtasks (such as counting the number of visits to each URL in a massive log). Each subtask processes a small portion of the data and does not depend on the intermediate results of other subtasks.
- An analogy: MPI is like a symphony orchestra. Hadoop is like a courier company's sorting center.

Example: P2P Download



Example: the ‘Slowest’ Distributed System



Search Alien Civilization, using volunteer PCs on the internet

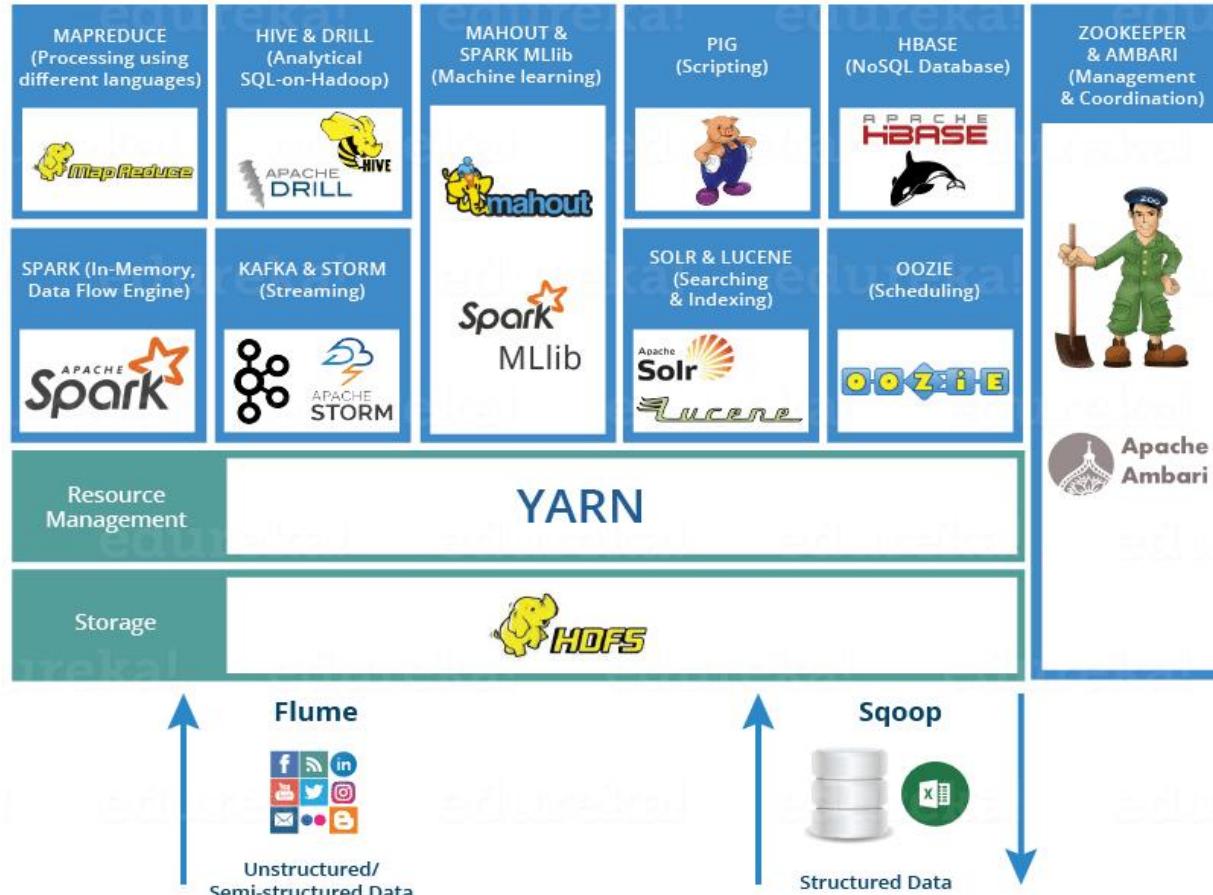
- Found nothing...
- Fermi Paradox



What is Hadoop?

- An open-source, reliable, and scalable distributed computing framework developed by the Apache Foundation.
- Core ideas: Derived from Google's GFS and MapReduce papers.
- Key advantages: High fault tolerance, high throughput, and low cost (using commodity hardware).

Hadoop Ecosystem



Hadoop Architecture

Hadoop

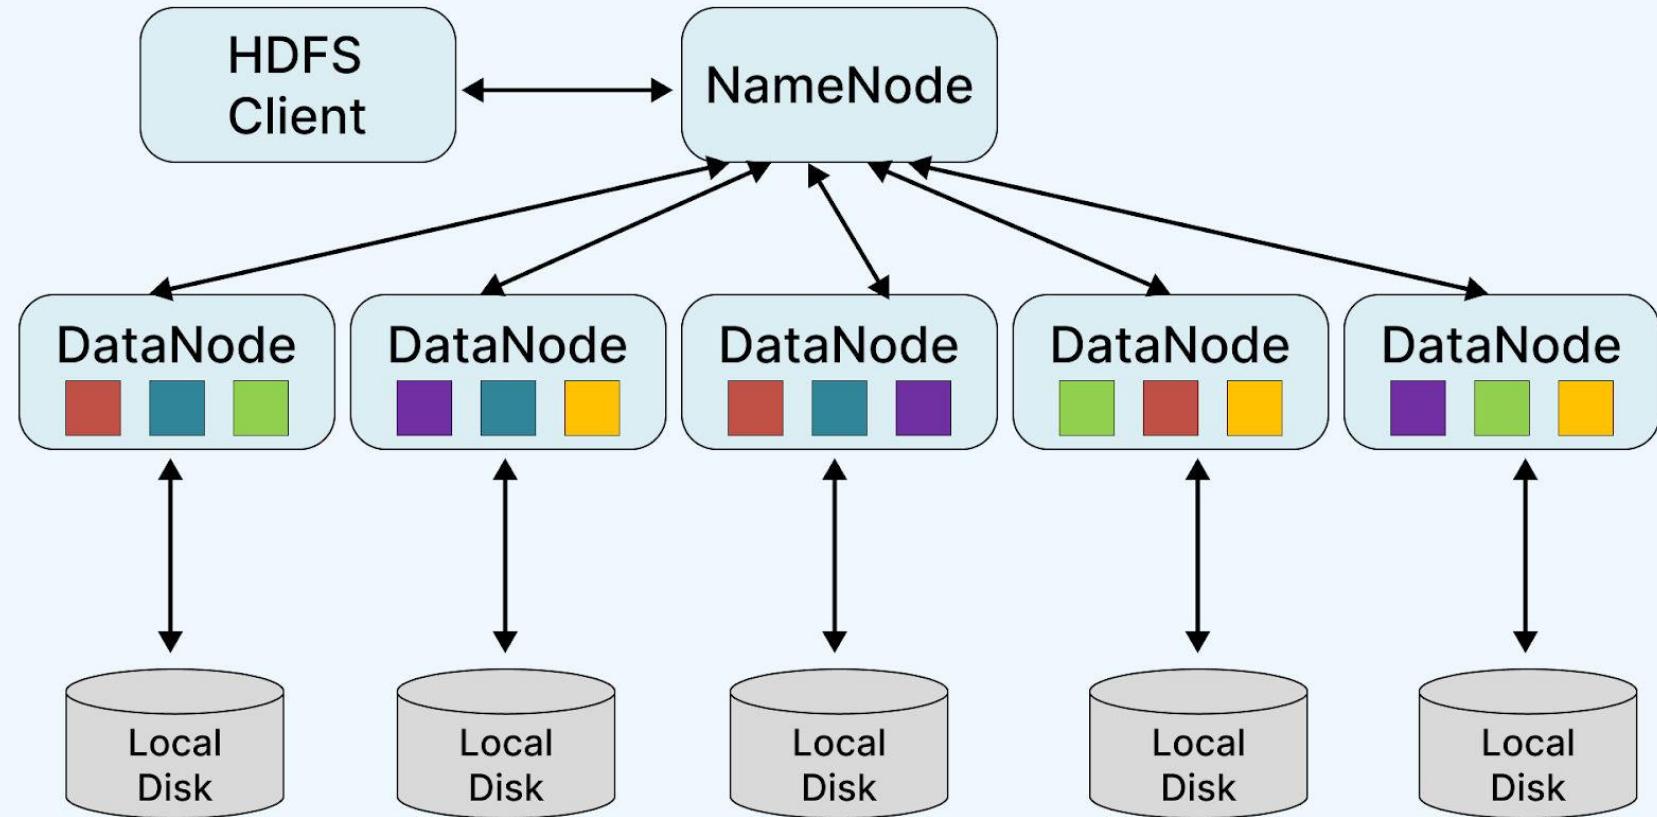
MapReduce
(Distributed Computation)

HDFS
(Distributed Storage)

YARN Framework

Common Utilities

Hadoop Cluster

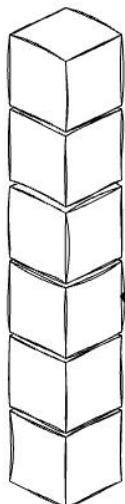


What is MapReduce?

- A programming model and computing framework for large-scale data processing.
 - The core concept is "divide and conquer."
 - Break down complex tasks into multiple smaller tasks, process them in parallel across a cluster, and then aggregate the results.

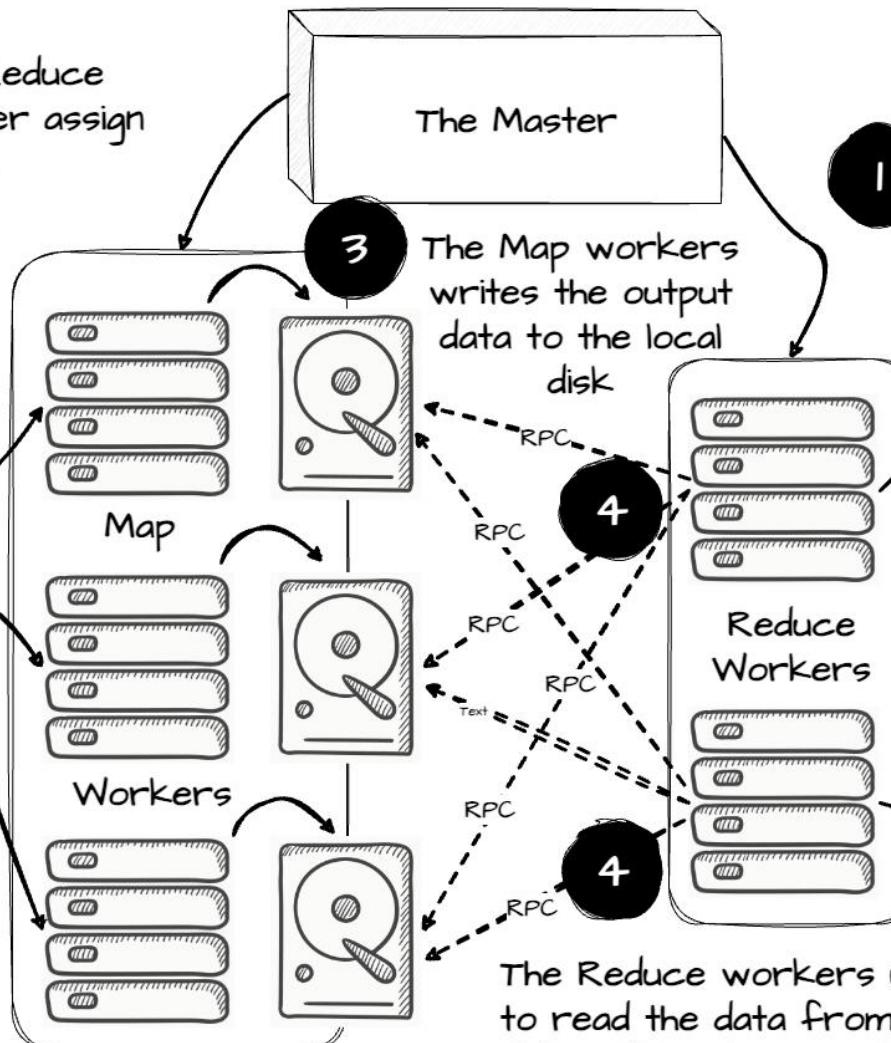
I

The MapReduce Master assign Map....



2

Map workers reads the input file splits.



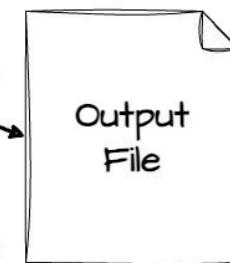
....and Reduce tasks to idle workers

I



5

The output of the Reduce function is appended to a final associated output file



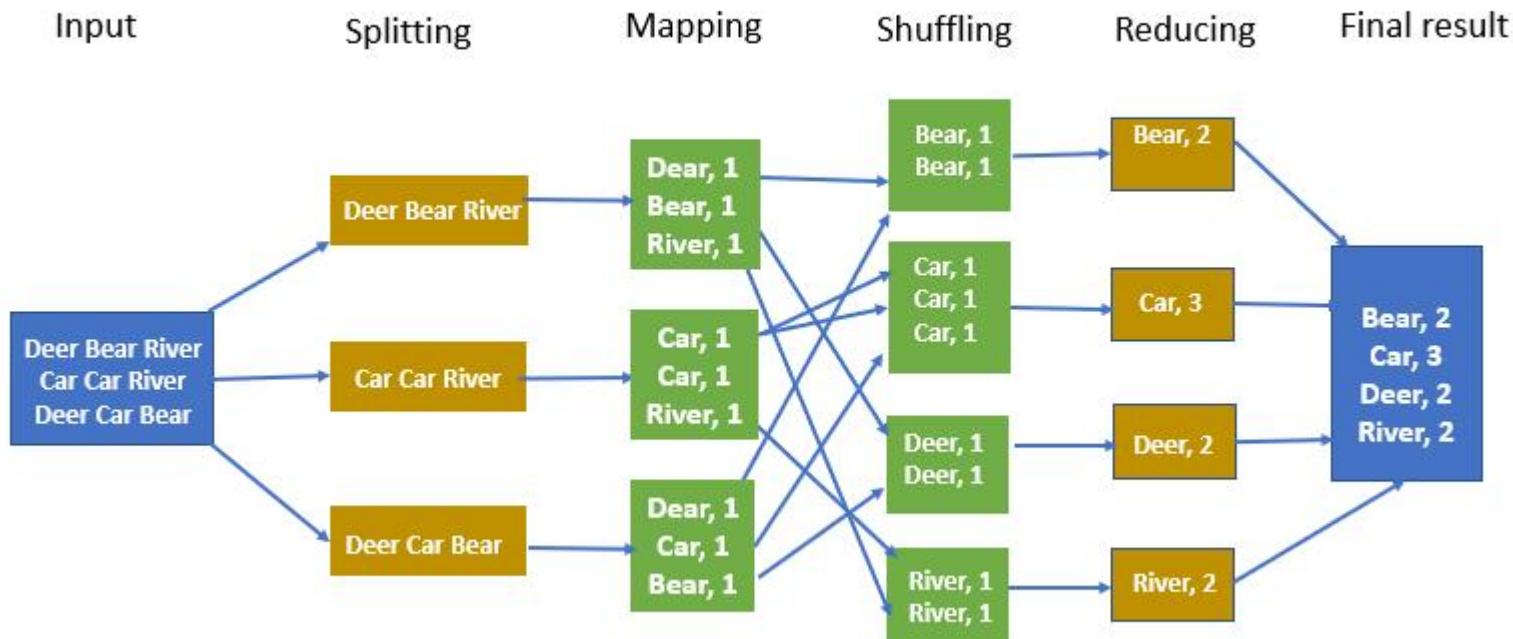
5

The Reduce workers use RPC to read the data from the local disks of the Map workers.

A vivid example: counting word frequencies in a library

- Question: How can we quickly count the number of occurrences of each word in all the books in the library?
- Traditional method: One person counts each book one by one. Too slow!
- MapReduce Method:
 - Step 1: Map: Have multiple people (workers) each be assigned a few books, and each count the number of occurrences of a word in their own book, outputting an intermediate result such as (word, 1).
 - Step 2: Shuffle & Sort: Group and sort all intermediate results by word, so that all records for the same word are grouped together.
 - Step 3: Reduce: Have multiple people (workers) each be responsible for a subset of words, and sum up all the counts for each word to produce the final result (word, total counts).

The overall MapReduce word count process



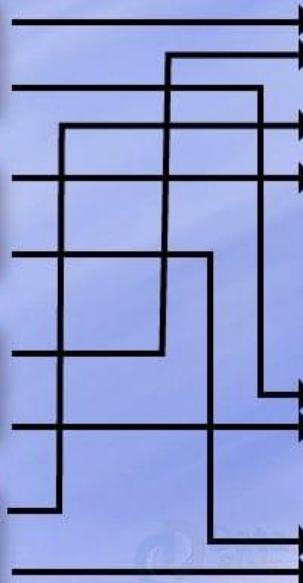
Code

- <https://www.geeksforgeeks.org/python/hadoop-streaming-using-python-word-count-problem/>

Shuffling & Sorting in Hadoop

**Output
From
Mapper**

Ayush	432
Mona	467
Bittu	898
Disha	436
Disha	978
Ayush	345
Bretty	456
Mayank	967



**Input
to
Reducer**

Ayush	432
Ayush	345
Bittu	898
Bretty	456
Disha	436
Disha	978
Mona	467
Mayank	967

- MapReduce = Map + Shuffle + Reduce